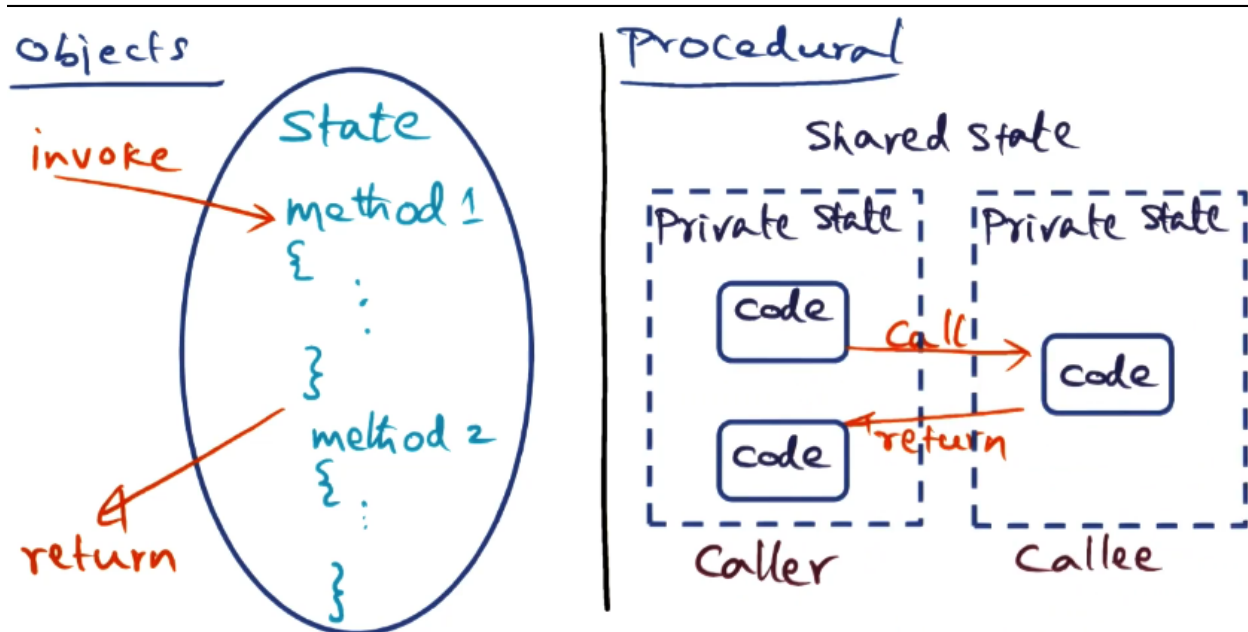


# Distributed Objects and Middleware

## Spring Operating System

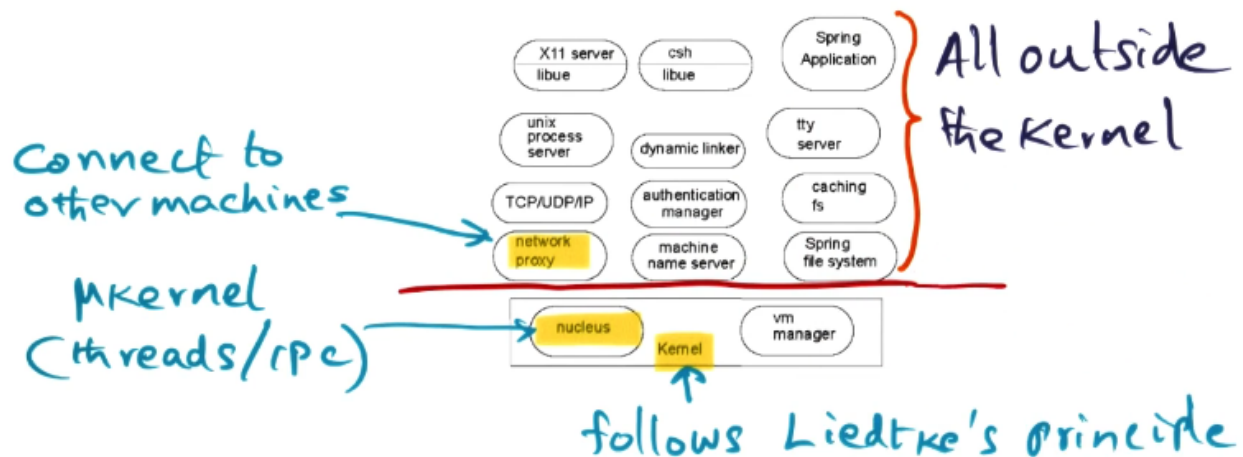
- \* Spring Operating System Introduction
  - Spring became Sun's Solaris operating system
  - Created by Yousef Khalidi as a grad student at Georgia Tech
    - + "Clouds" distributed operating system
- \* How to Innovate OS
  - Brand new OS or better implementation of known OS?
  - Marketplace needs
    - + Large complex server software
    - + Legacy applications running on existing OS, so industry tends to innovate rather than reinvent
  - Intel inside: External interfaces stay the same, but there is innovation in the microarchitecture
  - Sun took this approach ("Unix inside")
    - + Externally looked like Unix, but provide new APIs for new features
    - + Preserve everything good in standard OS, but also provide the capability for extensibility and flexibility
- \* Object based vs Procedural Design
  - Procedural: Some shared state, some private state. Subsystems interact through procedure calls
    - + State is all over the place (monolithic systems)
  - Object-based: Objects contain the state entirely within the object
    - + Methods within the object to manipulate the state of the object
    - + Strong interfaces that completely isolate discrete objects
    - + Performance consideration: Border crossings across logical protection domains can be slow
    - + Tornado also took an object-based approach to building OS kernels



Object-based vs Procedural Design

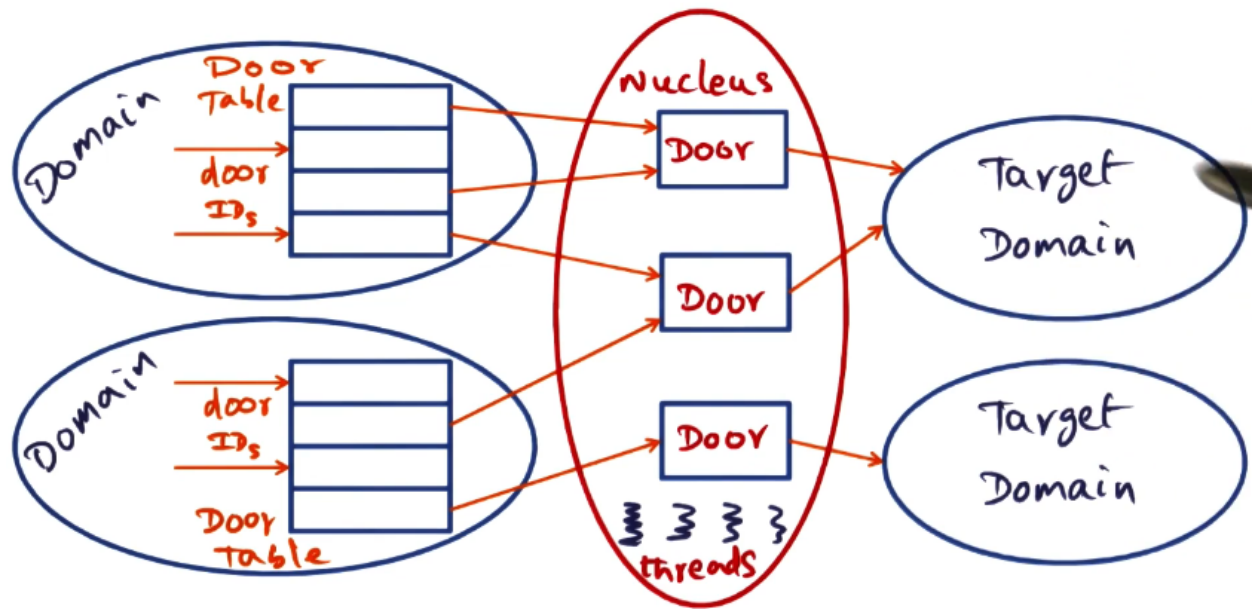
- \* Spring Approach

- Strong interfaces: Only thing exposed outside the subsystem is what services are provided by that subsystem, but not how
  - + "How" can be changed at any given time
  - + Leads to object-orientation
  - + Open, flexible, and extensible
  - + Don't want to be tied to a specific language
  - + Sun used Interface Definition Language (IDL) to expose interfaces
- Extensibility leads to microkernel-based approach
  - + Nucleus provides abstractions of threads and IPC
  - + Virtual memory manager provides memory management
  - + Microkernel is composed of these two systems
  - + Follows Liedtke's design principles
- Network proxy is interface that allows connection across machines



## Spring Design

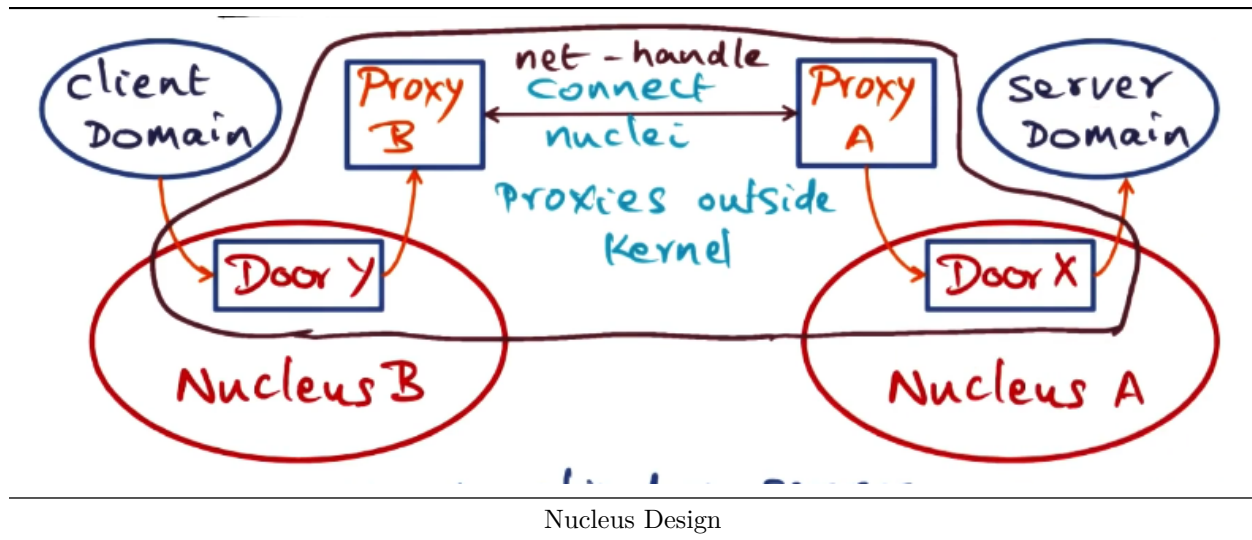
- \* Nucleus Microkernel of Spring
  - "Microkernel" of Spring
  - Only manages threads and IPC
  - Domain: Similar to a Unix process (address space)
  - Door: Provides access to a domain (entry points for target domain)
    - + Analogy: Using `fopen` provides a file descriptor to access a file
    - + Same idea, but for processes
  - Nucleus is involved in every door call
    - + Nucleus validates that the domain has the permission to use the door handle
    - + Client thread is deactivated and the thread is allocated to the target domain so it can execute the invocation
    - + Client thread is reactivated once the target domain is complete
    - + This makes cross-domain calls very fast
  - Kernel is composition of Nucleus plus memory manager



Nucleus Design

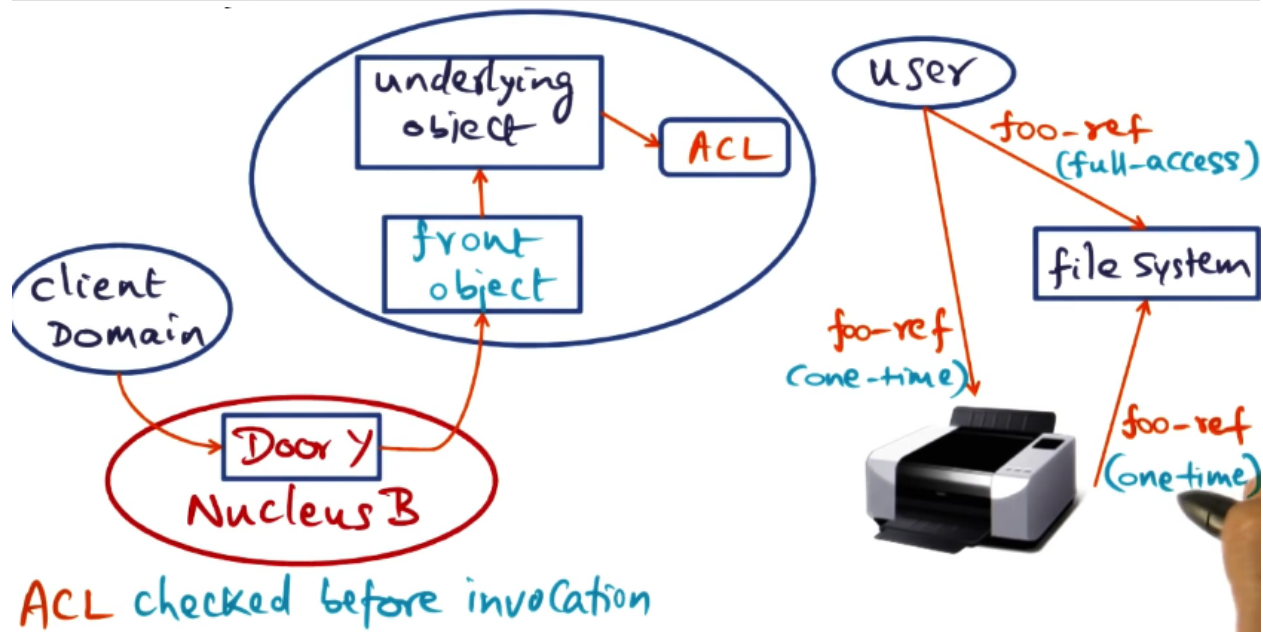
#### \* Object Invocation Across the Network

- Object invocation across the network is extended using network proxies
  - + Proxies can potentially employ different protocols (LAN, WAN, etc)
- Key property of building network operating system
- Proxies are invisible to client and server (unaware if they are on the same machine or different machines)
- Steps for invoking an object across the network
  1. Instantiate a proxy on the server node and establish a door for communication between proxy A and the server domain
  2. Proxy A exports a network handle embedding door X to Nucleus A
  3. Proxy B has door Y to connect to Nucleus B
  4. Proxy B uses net handle to connect to nuclei
  5. When the client wants to make an invocation on the server domain, it accesses proxy B which communicates with proxy A with door Y which communicates with the server domain through door X
- Proxies are outside the kernel



\* Secure Object Invocation

- Server object might need to provide different privilege levels for different clients
  - + File server might provide different access protections
- Spring provides this through a "front object"
- Client domain can only access front object which arbitrates access to the underlying object through the access control list (ACL)
- Policies for accessing underlying object can be implemented through different instances of the front object
- Client can also limit access to its objects
  - + Client has full access to file system
  - + Client passes a one-time reference to the printing object, which passes it to the file system
- Network accesses are fast and secure due to this implementation
- External Unix interfaces are identical, but Spring has modified the underlying implementation using object technology



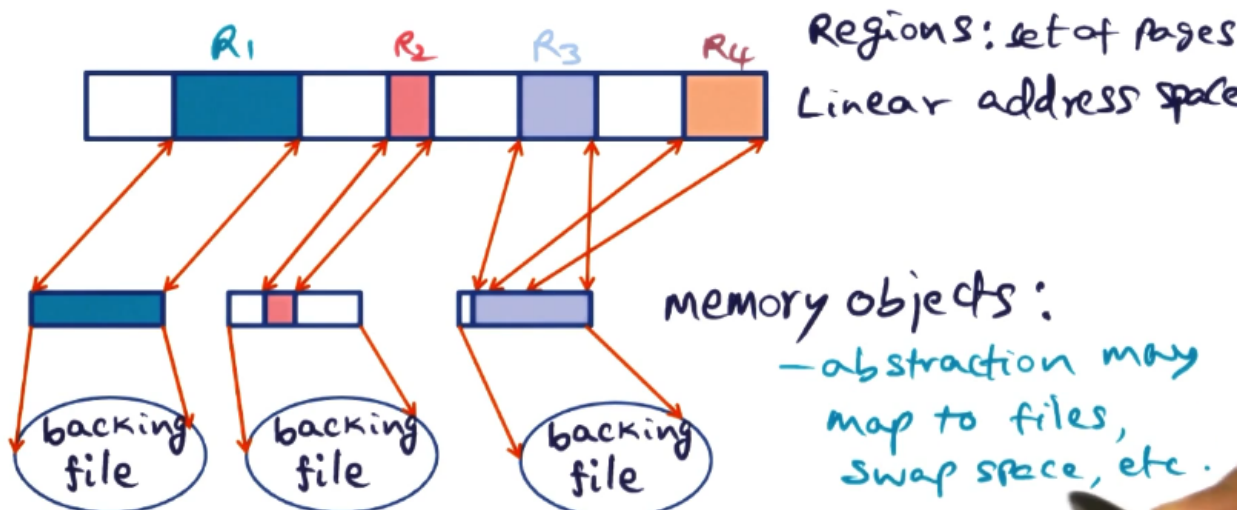
### Secure Object Invocation

#### \* Abstractions

| Feature   | Nucleus | Liedtke |
|-----------|---------|---------|
| Threads   | X       | X       |
| IPC       | X       | X       |
| Add Space |         | X       |

#### \* Virtual Memory Management in Spring

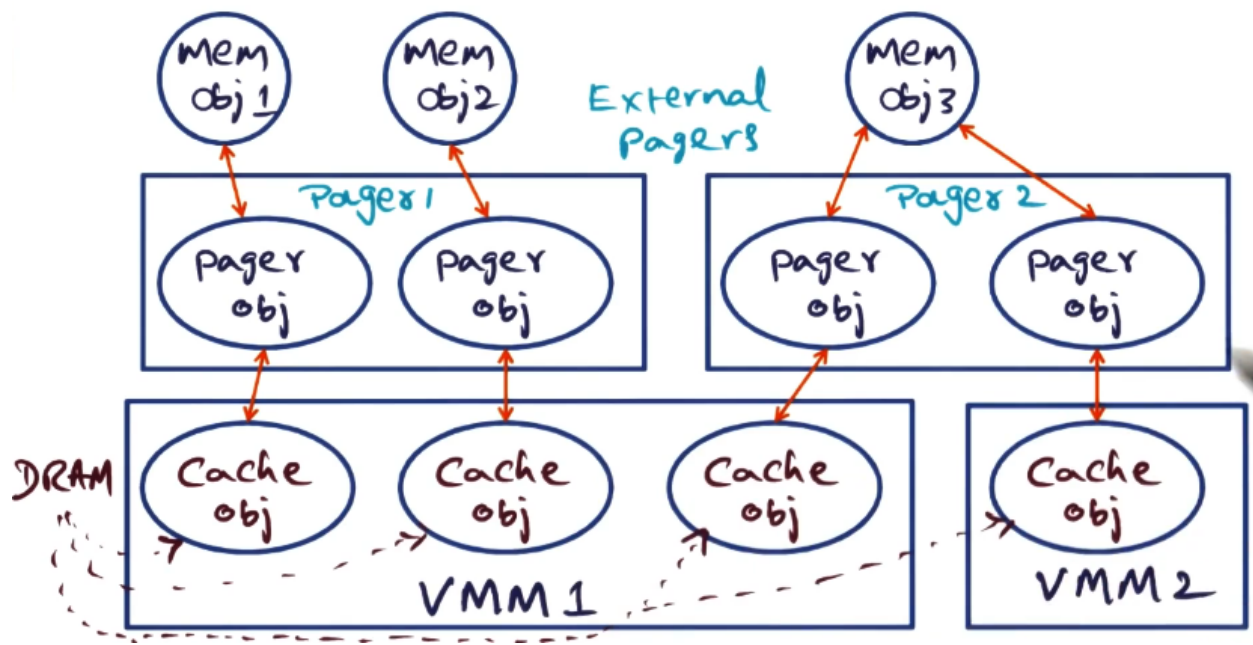
- Linear address space broken into regions (sets of pages)
- Memory objects: Abstraction that may map a portion of the address space to files, swap space, etc. (backing entities)
  - + Multiple memory objects can map to the same backing entity



Memory Management in Spring

#### \* Memory Object Specific Paging

- Memory object: Virtual memory representation
- Pager object: Establishes connection between virtual memory and physical memory
  - + Creates a cached objection representation in DRAM with the virtual memory management object
  - + Address space manager can make any number of mappings
  - + No single paging mechanism used for all memory objects
- Cache object: Physical memory representation
- If two VMMs have mappings to the same memory object, it is the responsibility of the pager object to coordinate the coherence of the cached objects, if this is needed
- In a single linear address space, you can have multiple pager objects managing different regions of the same address space



Paging in Spring

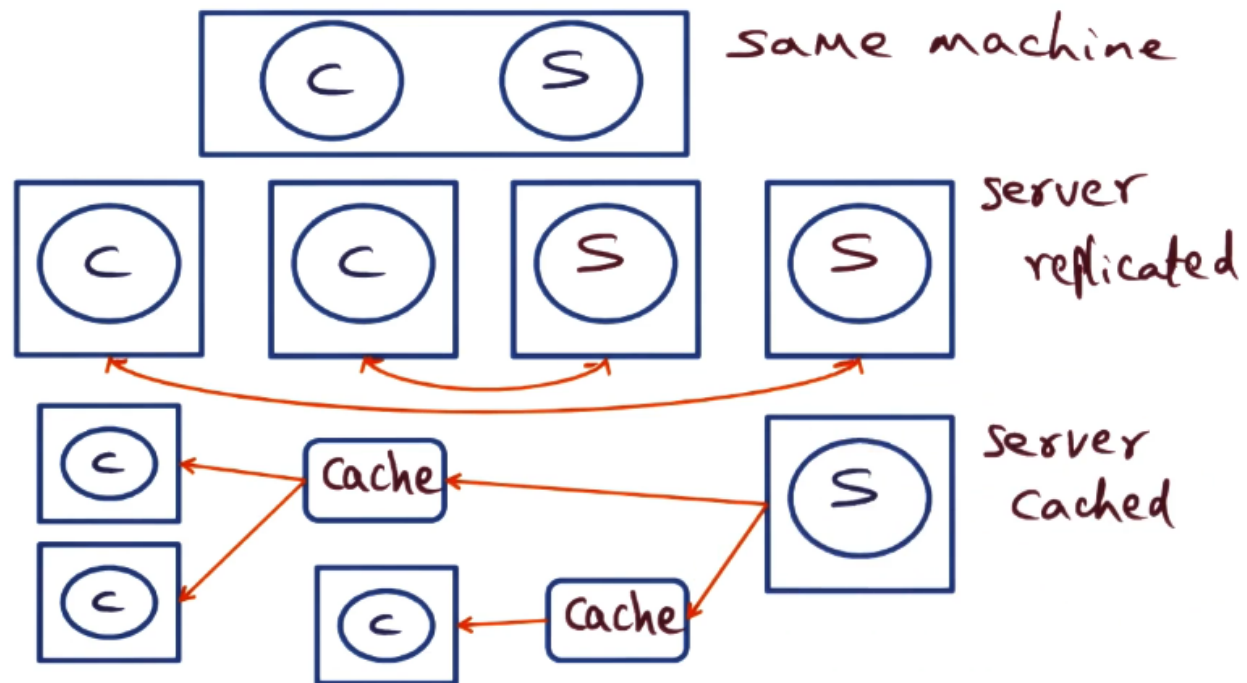
#### \* Spring System Summary

- Object oriented kernel
  - + Nucleus: Threads and IPC
  - + Microkernel: Nucleus and address space
  - + Door and Doot table: Basis for cross-domain calls
  - + Object invocation and cross machine calls
  - + Virtual memory management: Address space object, memory object, external pages, cached objects
- Tornado also used object technology, but Tornado used clustered objects as a representation for implementing kernel services
- In Spring, object technology permeates the entire design. It's a structuring mechanism as opposed to only an optimization

#### \* Dynamic Client Server Relationship

- In Spring, clients and servers of the network OS can be on the same machine or different machines
  - + The client/server interaction should be freed from their physical location
- Client requests are routed to different servers depending on physical proximity and current load
  - + Similar to how Google works today
- Server can be cached or replicated





Dynamic Client Server Relationship in Spring

#### \* Subcontract

- Subcontracts make dynamic relationship between client and server possible
  - + Analogous to offloading work to a third party
- Contract between client and server established through IDL
- Subcontract is the interface provided for realizing the IDL contract
- Client doesn't know or care if the server is a singleton, replicated
- Client side stub generation is simplified
- Subcontract responsible for details
- Can seamlessly add functionality to existing services using the subcontract interface
- Abstracts the underlying server implementation from the client

#### \* Subcontract Interface for Stubs

- Subcontract specifies if server is on same machine, different machine, different processor on the same machine, etc
- Three mechanisms
  1. Marshaling/unmarshaling for arguments: Subcontract marshals appropriately based on where the server is, abstracted from client
  2. On client
    - + Invoke: Calls function on server
  3. On server
    - + Create: Makes new service
    - + Revoke: Remove a service
    - + Process: Ready to process invocation requests

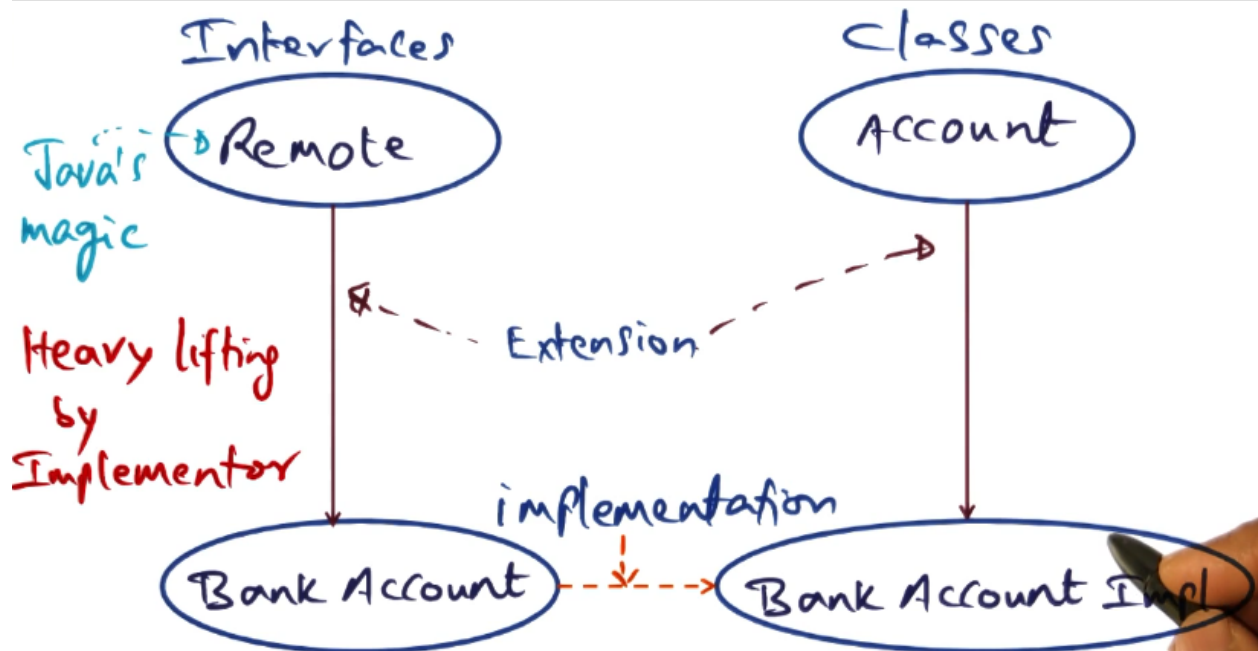
#### \* Spring Operating System Conclusion

- Sun was still building Unix boxes, but had completely revolutionized the structure of the network operating system through object technology
- This formed the basis for Java RMI



## Java RMI

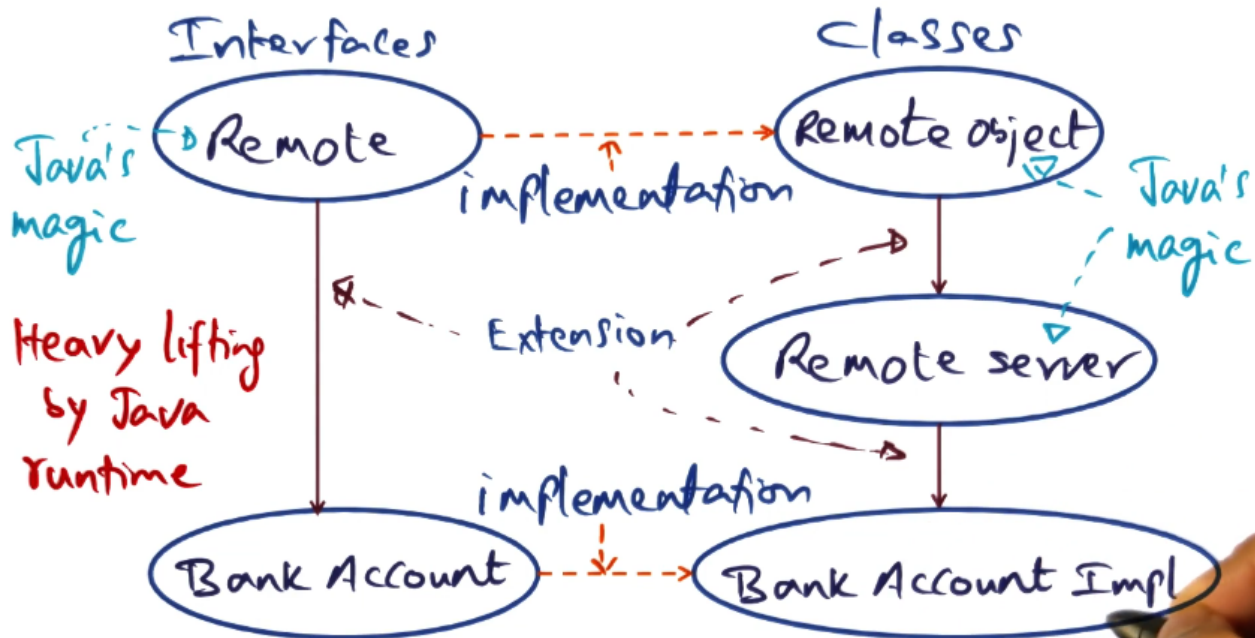
- \* Java History
  - Invented by James Gosling at Sun
  - Originally called Oak, later Java
  - Originally invented for use on embedded devices in the early 1990s
  - Intended for use with PDAs, then to set-top boxes, and then onto the Internet for powering E-commerce
- \* Java Distributed Object Model
  - Much of the heavy lifting involved in creating a client/server system with RPC (marshaling, unmarshaling, network interface) are subsumed by the Java distributed object runtime
  - Remote object: Objects that are accessible from different address spaces
  - Remote interface: Declarations for methods in a remote object
  - Failure semantics: Clients deal with RMI exceptions
  - Similarities/differences to local objects
    - + Object references can be parameters
    - + Parameters only passed as value/result
    - + In Java, an object can be passed to a method by reference and modified. A remote object cannot be.
- \* Bank Account Example
  - Server API
    - + Deposit
    - + Withdraw
    - + Balance
  - How best to implement using Java?
- \* Reuse of Local Implementation
  - Developer uses Java's remote interface to publish an existing local interface
  - The developer has to do the heavy lifting



Using Local Implementation of Bank Account

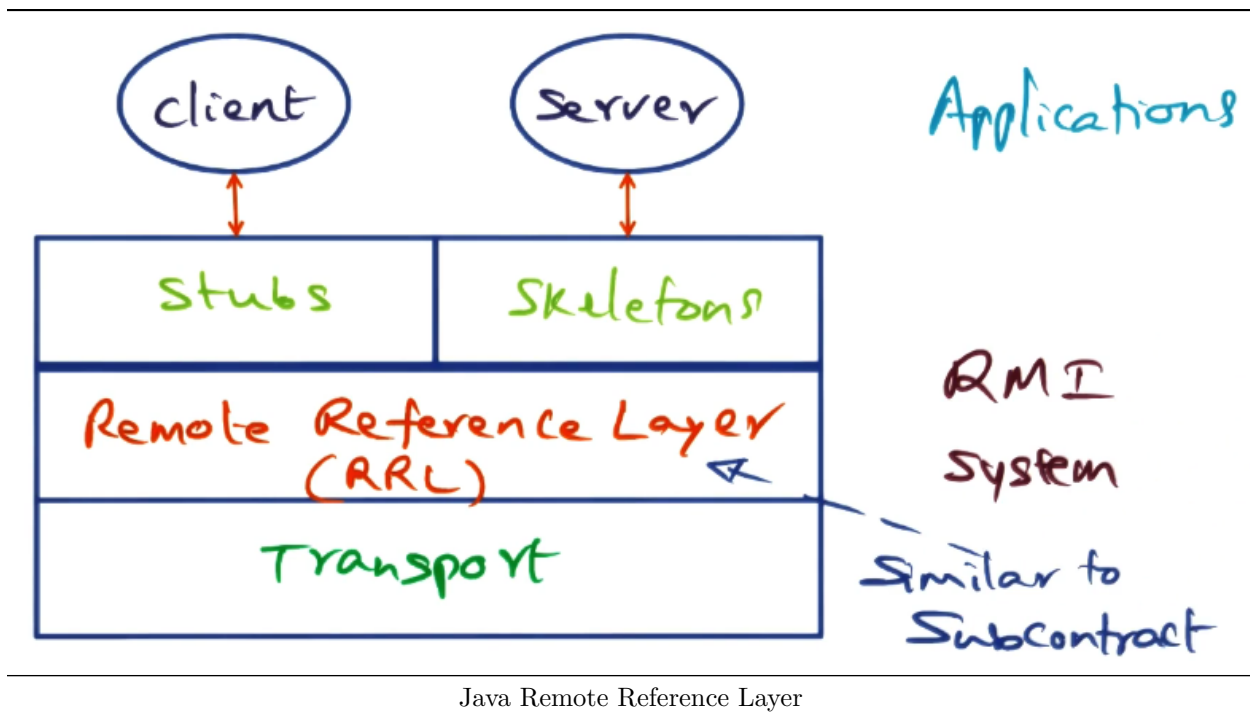
\* Reuse of Remote

- Developer writes the local object and uses the remote interface, identical to the local example
- The bank account implementation is derived from the Java built-in classes for remote object and remote server, not the Account superclass
- This results in the heavy lifting needed to make the bank account object visible to network clients is done by the Java runtime
- The remote implementation is preferable because the Java RMI is doing all of the heavy lifting



Using Remote Implementation of Bank Account

- \* Java RMI at Work (Server)
  - `BankAccount acct = new BankAccountImpl();`
  - `URL url = new URL("mywebaddress");`
  - `Java.rmi.Naming.bind(url, acct);`
- \* Java RMI at Work (Client)
  - `BankAccount acct = Java.rmi.Naming.lookup(url);`
  - `float balance;`
  - `acct.deposit($);`
  - `acct.withdraw($);`
  - `balance = acct.balance();`
  - Failures will throw remote exceptions back to the client
    - + Client will not know where exactly the server failed
- \* Java RMI Implementation (Remote Reference Layer)
  - Remote Reference Layer (RRL):
  - Client side stub initiates an RMI call using the RRL
    - + All of the marshaling/unmarshaling is handled within the RRL
  - Skeleton exists to unmarshal arguments from the client using the RRL
    - + Makes the call up to the server, marshals the result, and goes through the RRL to send it
  - There could be several instances of the server; the RRL handles where the server is, is it replicated, how it's handling requests, etc
  - Derived from the subcontract mechanism from Spring



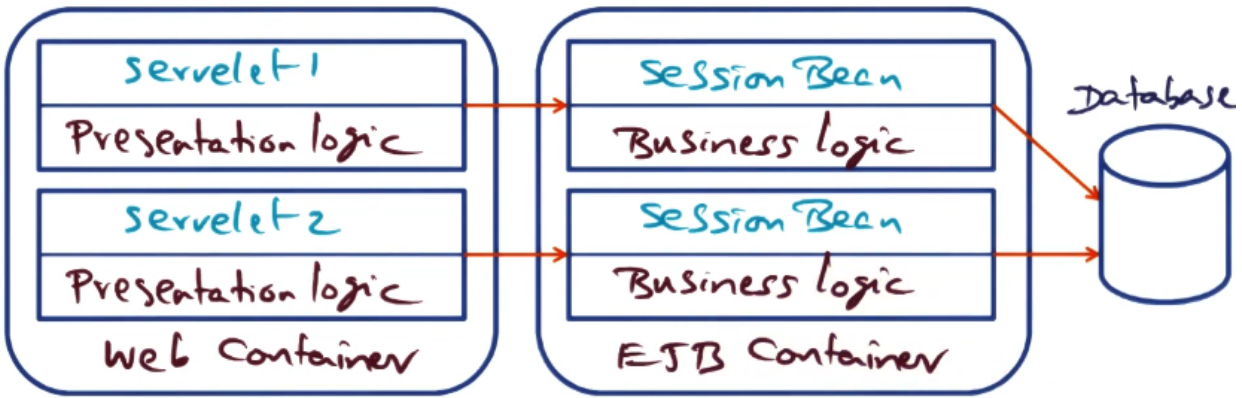
- \* RMI Implementation Transport
  - Provides the following connections:
    - + Endpoint
    - + Transport
    - + Channel
    - + Connection
  - Endpoint: Protection domain (Java VM) with a table of remote objects it can access
    - + Server/client can exist within this sandbox
  - Connection Management
    - + Setup, teardown, listen
    - + Liveness monitoring
    - + Choice of transport
  - Transport: Listens on a channel
    - + When an invocation arrives, the transport is responsible for identifying the dispatcher on the domain (endpoint)
  - Channel: Medium over which communication occurs
    - + TCP or UDP
  - RRL decides the right transport to use depending on the location of the two endpoints
- \* Java RMI Conclusion
  - When the time is right, ideas from research form the basis for useful tools in industry

## Enterprise Java Beans

- \* Enterprise Java Beans Introduction
  - How do we structure the software for a large scale distributed service?
  - Java bean: Many Java objects in a bundle that can be passed easily between applications
- \* Inter Enterprise View

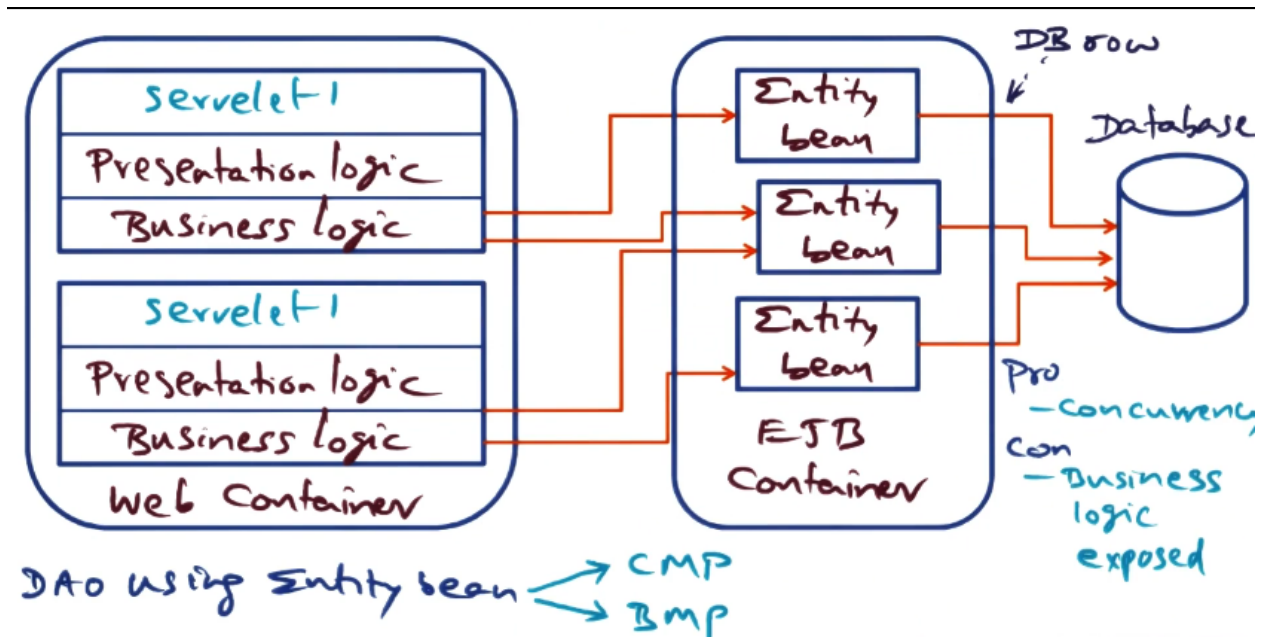
- Intra Enterprise View: Many services and servers interconnected
  - Inter Enterprise View: Service may not be serviced by a single entity.
- Instead, many enterprises work together
- + Supply chain model
- Enterprise challenges
    - + Interoperability
    - + Interface compatibility
    - + System evolution
    - + Scalability
    - + Reliability
- \* Enterprise Java Beans Example
    - Giant scale services: Airline reservation, Gmail, web surfing as opposed to a local file service
    - There are many common features across services, such as a shopping cart
    - Don't want to reinvent the service every time
      - + Object technology provides the ability to reuse components
  - \* N Tier Applications
    - Presentation layer: Paints the screen on browser
    - Application logic corresponding to what the service provides
    - Business logic corresponding to pricing, etc.
    - Database layer to access information for application and business logic
    - Issues in distributed programs:
      - + Persistence
      - + Transactions
      - + Caching
      - + Clustering
      - + Security
    - Reduce network communication (latency), reduce security risks, improve concurrency in responding to individual requests
      - + "Embarrassingly parallelizable" application
    - Desire to reuse components of application logic as much as possible
  - \* Structuring N Tier Applications
    - Using JEE framework for structuring applications
    - Container: Protection domain, typically in the form of a JVM
      1. Client: Interact with browser on end client
      2. Applet: Interact with browser on end client
      3. Web: Presentation logic
      4. EJB: Business logic (Enterprise Java Beans)
    - Bean: Unit of reuse (bundle of Java objects, such as shopping cart)
      1. Entity: Row of a database (persistent objects with primary keys)
        - + Persistence can be built into the bean or container
      2. Session: Associated with a client and temporal window (session)
        - + Could be stateful or stateless
      3. Message: Useful for asynchronous behavior
    - Containers host beans
    - Fine-grained beans provides greater opportunity for increasing concurrency of requests
      - + Makes business logic more complex
  - \* Design Alternative (Coarse-grain Session Beans)
    - One session bean responsible for the specific needs of the client it is servicing
      - + Handles accesses to database
      - + Multiple sessions within EJB depending on number of clients
      - + Business logic is confined to the corporate network since it's

- within the EJB container, not the web container
- Pros:
  - + Minimal container services
  - + Business logic not exposed
- Cons:
  - + App structure akin to monolithic kernel



Coarse-grain Session Beans

- \* Design Alternative (Data Access Object)
  - Want parallelism in accessing database that coarse-grained doesn't provide
    - + Database accesses are slow due to disk and network speeds
    - + Reduce latency resulting from this slowness
  - Presentation and business logic are contained within the client applet
  - Database accesses are handled by entity beans in EJB container
    - + Entity bean can be one or many rows of database
    - + Reduces time for database access by exploiting concurrency
    - + An entity bean could cluster requests and amortize database access if multiple requests are querying the same row
  - Persistence is required for data access object
    - + Bean managed persistence
    - + Container managed persistence
  - Pros
    - + Concurrency
  - Cons
    - + Business logic exposed

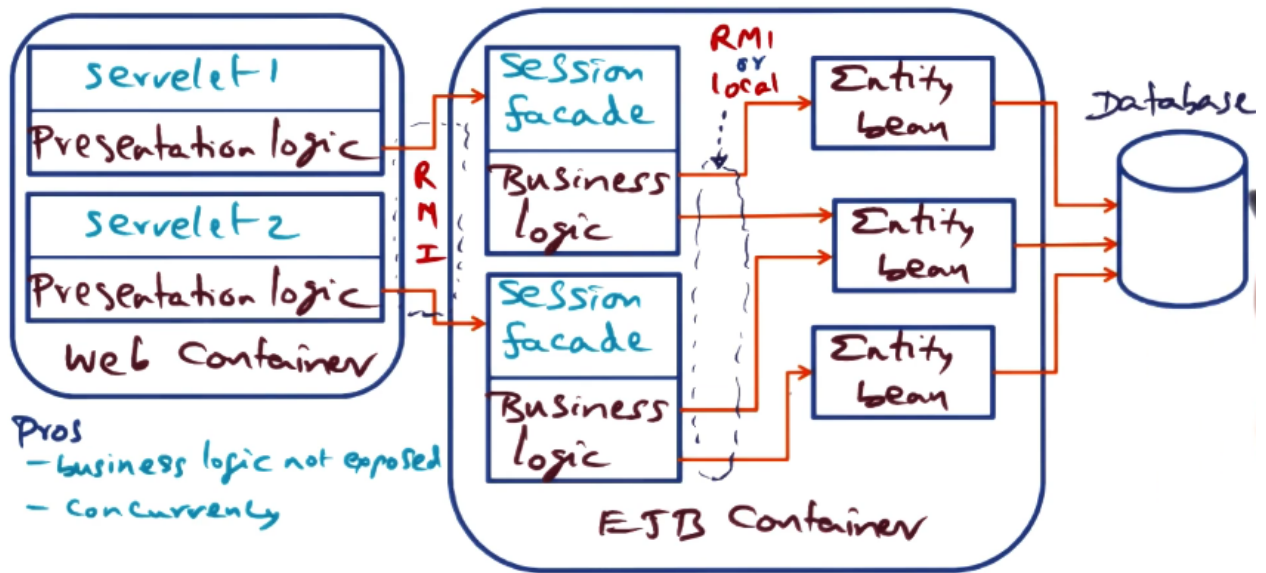


Data Access Object

\* Design Alternative (Session Bean with Entity Bean)

- Use both session and entity beans to achieve parallelism without exposing the business logic
- Session facade: Worries about data access needs of associated business logic
- Can farm parallel requests to multiple entity beans
- Data access objects still use entity beans
- Web container uses RMI to communicate with the business logic
- Session facade communicates with entity bean with RMI or locally
  - + Local communication eliminates the overhead of the network
- Pros
  - + Business logic not exposed
  - + Concurrency
- Cons
  - + Additional network access, but can be mitigated by colocating session facade with entity bean





Session and Entity Bean

- \* Enterprise Java Beans Conclusion
  - EJB allows developers to develop business logic without being concerned with security and network interfaces