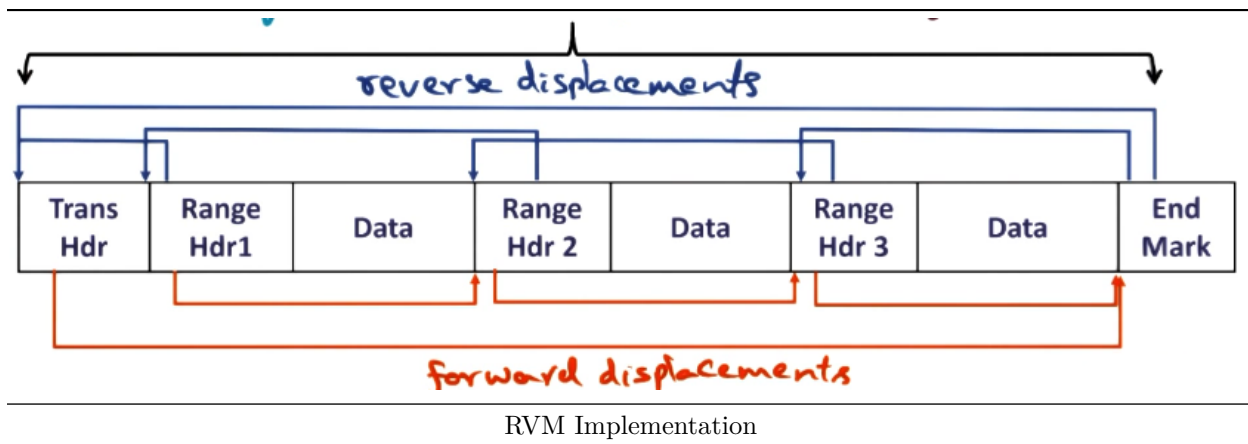


# Failures and Recovery

## Lightweight Recoverable Virtual Memory

1. Lightweight Recoverable Virtual Memory Introduction
  - How do we build systems that survive crashes?
    - Hardware, software, power failures
  - LRVM: Persistent memory layer in support of system services
2. Persistence
  - Why?
    - Need of OS subsystems (filesystem inodes)
    - Log memory to disk, but need to keep it consistent with changes in memory
  - How?
    - Make virtual memory persistent
    - All data structures contained in virtual memory become persistent
    - Subsystems don't have to worry about flushing to disk if it's handled by the OS
  - Who will use it?
    - Subsystem designers if it is performant
  - How to make it efficient?
    - Writing every change in memory to disk incurs significant overhead due to the latency of writing to disk (seek, rotational)
    - Use persistent logs to record changes to VM (similar to xFS)
    - Only write logs of the changes, buffer to save them all at once to solve the small write problem
3. Server Design
  - Distinguish between persistent metadata and normal data structures
  - Create external data segments to back persistent data structures
    - Applications manage their persistence needs
  - Designers can make regions of virtual memory to segments on disk
    - Designer's choice to use single or multiple data segments
  - Mapping from virtual address to external data segments is 1-to-1
    - No overlap in occupancy of virtual address space
    - Simplifies design
  - Application can map and unmap whenever needed
4. Recoverable Virtual Memory Primitives
  - Initialization
    - `initialize(options)`: Log segment to be used for persistent data
    - `map(region, options)`: Virtual address to external data segments
    - `unmap(region)`: Deletes a mapping
  - Body of server code (like a critical section)
    - `begin_xact(tid, restore_mode)`: Start of changes to log
    - `set_range(tid, addr, size)`: Only using a portion of address range
    - `end_xact(tid, commit_mode)`: Log changes to log segment
    - `abort_xact(tid)`: Discard changes to persistent data structures
  - GC to reduce log space (done by LRVM automatically)
    - `flush()`: Provided for application flexibility
    - `truncate()`: Provided for application flexibility
  - Miscellaneous
    - `query_options(region)`
    - `set_options(options)`
    - `create_log(options, len, mode)`
  - When a developer is making changes within a transaction, RVM commits the changes to a redo log in the log segment
    - These are only committed to disk if the transaction isn't aborted
    - Committing to disk happens at opportune times

- Truncation: Throwing away the redo log after it's committed
  - Developer can explicitly manage its redo log as a way to optimize the use of disk space
- Small set of primitives that are easy to use and performant
- Transaction: Intended for recovery management, doesn't need all of the properties associated with typical database transactions (ACID)
  - Atomicity
  - Consistency
  - Isolation
  - Durability
- RVM doesn't allow for nested transactions or support concurrency control
  - Developer must implement at a higher level if needed
- 5. How the Server Uses the Primitives
  - Initialize address space from external segments
  - `begin_xact(tid, mode);`
    - `set_range(tid, base_addr, #bytes);`
    - write metadata m1; // contained in range
    - write metadata m2; // contained in range
  - `end_xact(tid, mode);` // can also be abort
  - When developer calls `set_range`, LRVM creates an undo record
    - Copy of virtual address space for the specified number of bytes
  - Mode specifier allows user to specify to RVM whether transaction will ever abort
    - If developer is certain transaction won't abort, can specify a `no_restore` mode so RVM knows not to create an undo record
  - When the application is writing metadata, no action is needed by LRVM
    - Changes happen directly to the virtual address space of that particular process where the in-memory copy of the persistent data structures are living
- 6. Transaction Optimizations
  - No-restore mode in `begin_xact`
    - No need to create in-memory undo record
  - No-flush mode in `end_xact`
    - No need to do synchronous flush redo log to disk
    - Lazy persistence (will be persistent, but not at the point of `end_xact`)
    - There's a period of vulnerability between `end_xact` and flush, but this provides improved speed by removing synchronous I/O
    - "Transactional systems perform well when you don't use transactions"
  - Use transactions as insurance
- 7. Implementation
  - Lightweight in terms of transactional properties
  - No undo/redo value logging
    - Undo log: Creating an undo record of changes only in memory, not on disk (only for duration of transaction)
    - Redo log: All changes to different regions between `begin_xact` and `end_xact` (also in memory, then flushed to disk)
  - On commit: Replace old value record in virtual memory with new value records (automatic based on how RVM works)
    - Must undo changes if a transaction aborts
  - Creating the undo log is optional if the developer guarantees that the transaction won't abort
  - Writing the redo log to disk can be done lazily (don't block on call to `end_xact`)
    - Gives a window of vulnerability where data can be lost
  - Can traverse redo log in both directions to provide flexibility



## 8. Crash Recovery

- Redo log has transaction header
  - Between transaction header and end mark contains all of the changes that have been made by a critical section
- Resume from crash
  - Read redo log from disk
  - Apply to external data segments in memory
- All of the needed information is contained in the redo log

## 9. Log Truncation

- If time between crashes is long (as we hope), log segments build up
  - Redo log will also build up
- When a crash happens, apply to redo log to disk to clear it
  - Clogging disk space, unnecessary overhead
- Truncating the log: Read logs from disk and apply to external data segments so the logs can be discarded
  - Simply apply crash recovery algorithm (same logic)
  - Read redo logs into memory and apply to data segment
- Log truncation: Perform in parallel with forward processing
  - LRVM splits log record into epochs
  - Truncation epoch: Crash recovery is using this
  - Current epoch: Server is using this
- Biggest challenge in LRVM is log truncation code
  - So much coordination involved

## 10. Lightweight Recoverable Virtual Memory Conclusion

- LRVM is classic systems research example
  - Understand pain point and create a solution
- Pain point: Managing persistence of critical data structures
- Solution: Lightweight transactions without typical ACID properties

## RioVista


### 1. RioVista Introduction

- Synchronous I/O makes transactions heavyweight in LRVM even though the semantics of transactions have been simplified considerably
- RioVista's goal is performance-conscious design of persistent memory

### 2. System Crash

- Two problems concerning failure
  - Power failure: Can we throw some hardware at the problem and make it disappear? (UPS)
  - Software crash: Reserve a portion of main memory that survives crashes

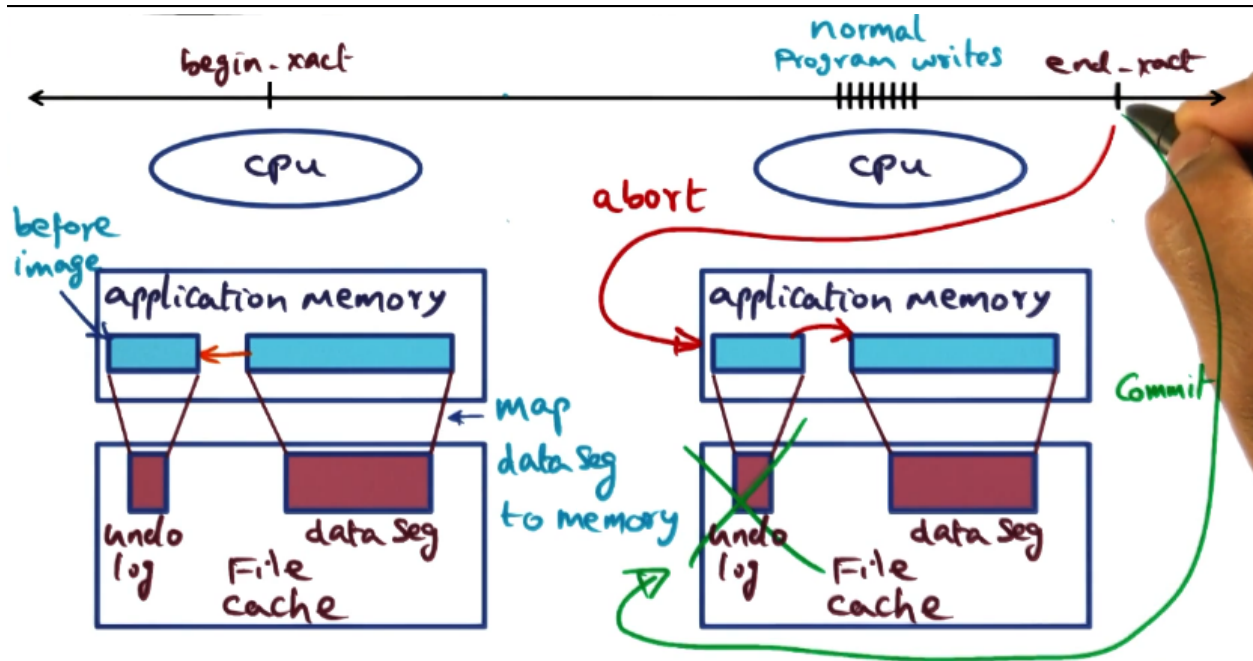
- RioVista is only concerned with software crashes
    - Makes transactions cheaper
3. LRVM Revisited
- begin\_xact: Memory copy of portion of memory by LRVM
  - body: Normal program writes
  - end\_xact: Disk copy by LRVM, get rid of undo logs
  - log truncation: Redo logs -> data, get rid of redo logs by LRVM
  - Upshot: 3 copies by LRVM to persist data
    - Biggest vulnerability is power failure if using no-flush option
4. Rio File Cache
- Rio: Using battery-backed DRAM to implement a persistent file cache
    - File system uses file cache holds data brought from disk
    - Persistent: Use UPS to persist file cache data
  - OS buffers writes to DRAM and writes to disk at opportune times
    - User has to use fsync to force writes to disk
  - mmap: Normal application memory becomes persistent
    - Backed by battery
  - Using battery-backed DRAM means no synchronous writes are required



Rio File Cache

5. Vista RVM on Top of Rio

- Vista: RVM library on top of Rio
- Semantics are identical to LRVM semantics, but the implementation takes advantage of the fact that it's sitting on top of Rio
- When mapping virtual memory to a data segment, it's mapped to the file cache which is already guaranteed to be persistent
- Specify address range to be persistent at the point of `begin_xact`
  - Create an in-memory copy of the memory to be modified, which serves as the undo log
  - Undo log is mapped to file cache (persistent by definition)
  - Changes during the body of the critical section are already persistent
- At `end_xact`, changes are committed
  - In Vista, changes are already committed because they're in persistent memory
  - Faster because there's no need to write to disk synchronously
- If the transaction aborts, the undo record created at the beginning of the transaction is simply copied back to memory
- Rio+Vista is fast because there's no disk I/O
  - Everything is memory resident and can be written back to disk asynchronously



Vista Recoverable Virtual Memory

6. Crash Recovery
  - Treat like an abort
    - Recover old image from undo log (survives crashes since it's in Rio file cache)
  - Crash during crash recovery?
    - Idempotency of recovery, so there's no issue
7. Vista Simplicity
  - 700 lines of code in Vista compared to 10000 in LRVM
    - Performs three orders of magnitude better than LRVM (no disk I/O)
  - Simpler
    - No redo logs or truncation code
    - Checkpointing and recovery code is simplified
    - No group commit optimizations
  - Upshot: Simple like LRVM but performance efficient
8. RioVista Conclusion

- Shows that by changing the starting assumptions of the problem, you can arrive at a completely different solution

## Quicksilver

### 1. Quicksilver Introduction

- Making recovery a first class citizen in an OS design, not an afterthought
- Performance and reliability are generally considered to be opposing concerns
- Quicksilver's approach is that if recovery is taken seriously at the initial design, you can have both

### 2. Cleaning up State Orphan Processes


- "Not responding" windows come from programs not being hygienic
- When an application closes, it needs to clean up any resources it was using
- LRVM and RioVista only target state that needs to be persisted across system crashes
- NFS is stateless; server doesn't maintain any state pertaining to the clients
  - The server cannot know about what state the client is reliant on when a transaction aborts. This state may live forever if a client application crashes

### 3. Quicksilver

- Quicksilver identifies many problems that we face with our everyday computing with orphan windows and such
- Built by IBM in the early 80s

### 4. Distributed System Structure

- Applications interact with system services, which interact directly with the microkernel
- Structure provides extensibility and high performance

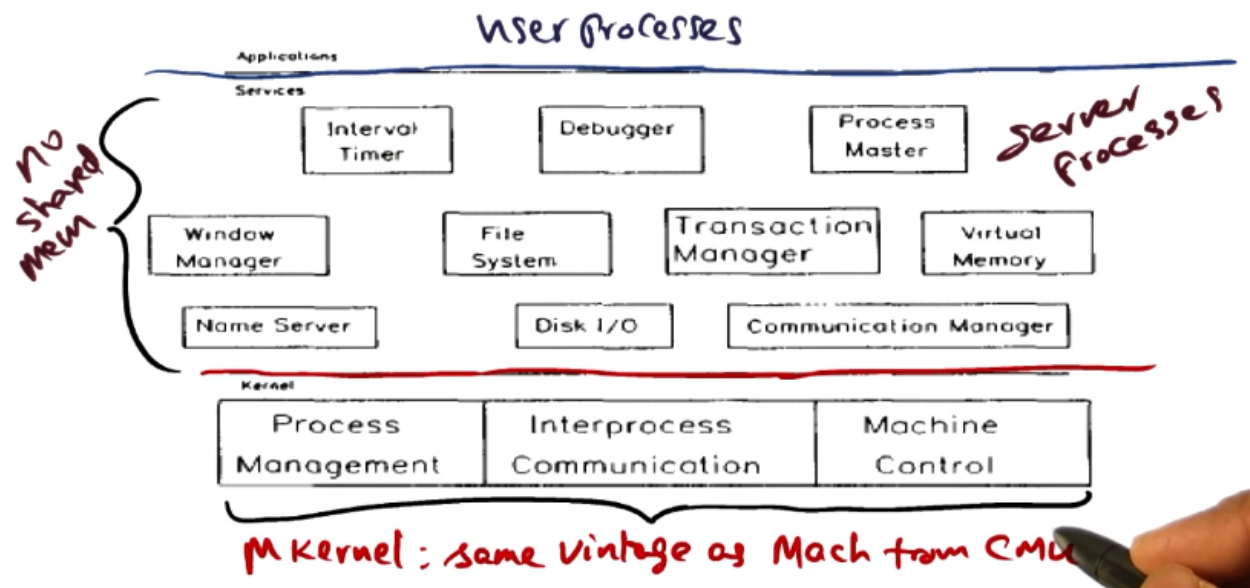


Distributed System Structure

### 5. Quicksilver System Architecture

- Quicksilver was architected to be very similar to modern day designs

- No shared memory
- Server processes all sit above microkernel
- Microkernel; same vintage as Mach from CMU
- 80s were moving from CRTs and mainframe to desktop computers
- Quicksilver proposes transaction as a unifying concept for recovery management of the servers
  - Transaction manager is service provided by OS to manage this



Quicksilver Architecture

#### 6. IPC Fundamental to System Services

- Service-q used to handle IPC; client adds a request to service-q and initiates an upcall to the server. Server executes request, completion goes back to service-q, and OS gives control back to client
  - Created by server
  - Any process can connect and any server can service requests
  - No loss of requests or duplicate requests
  - Service-q is globally unique, so client doesn't need to know where in the network the requests is being serviced
  - Supports synchronous and asynchronous requests
  - Very similar to RPC (invented around the same time)
- Recovery mechanism is intimately tied to RPC
  - Client/server interactions must use IPC
  - Binds recovery with IPC to make it cheaper
- Quicksilver created in early 1980s, paper not published until 1988
  - Industry only published when the system was finished
  - Nowadays, everybody publishes often



#### 7. Bundling Distributed IPC and Transactions

- Transaction: Secret sauce for recovery management
  - Lightweight, not heavyweight notion associated with databases
  - LRVM took their ideas for transactions from Quicksilver
- IPC calls are tagged with transaction headers
  - Want all failures to be recoverable, clean up state from clients, servers, communication manager, etc.
- Communication manager: One per node, handles IPC
  - Transaction data is piggybacked on top of regular IPC so there's no additional overhead




- Transaction manager: One per node, handles tracking local changes to commit/abort as part of a transaction later
  - Creator of transaction is default owner of transaction, coordinator for transaction tree that is being established (might go through several calls to other servers)
- Transactions need to clean up windows started by a window manager, file pointers opened by the filesystem, any additional state
  - Transaction manager needs to track this state across an entire transaction tree

Transaction  $\Rightarrow$  secret sauce for recovery management



Examples



#### Quicksilver IPC and Transactions

#### 8. Transaction Management

- Coordinator can be different from owner
- Client makes a call to a filesystem, filesystem makes call to data server
  - Client is owner of transaction tree, filesystem and data server are participants
  - Owner can delegate ownership to someone else
  - Clients are most fickle; clean up will become difficult if client crashes, so it might delegate to another node
- Heavy lifting done by Quicksilver is facilitating recovery management through the transaction tree

#### 9. Distributed Transaction

- Transactions are inherently distributed
- Results in a graph structure of the transaction tree
- Different types of failures
  - Client crashes
  - Connection failure
  - Subordinate transaction manager failed to report
- Transaction manager logs periodically to a persistent store to create checkpoint records for the client
  - Useful for recovery of work
- When a transaction manager terminates a transaction, all of the nodes must clean up all of their state

#### 10. Commit Initiated by Coordinator

- When a transaction completes, the transaction tree is traversed to clean up any resources created to satisfy the request



- Initiated by coordinator for commit or abort
  - Down the tree (initiated by coordinator)
    - Vote request
    - Abort request
    - End commit/abort
  - Up the tree
    - Response commensurate with request
  - If the transaction tree is representing the client/server relationship for opening a window...
    - If client crashes, coordinator will send an abort. Window manager will clear window since it's volatile
  - If the transaction tree is representing the client/server relationship a writing to a file...
    - If client crases, coordinator will send an abort. File system will undo changes to disk using checkpoints
11. Upshot of Bundling IPC and Recovery
- Reclaim resources: Service can safely collect all breadcrumbs left behind by failed clients and servers
    - Have a tree of all of the state changes to undo
  - No extra communication due to piggybacking off of IPC
  - Examples
    - Memory
    - File handles
    - Communication handles
    - Orphan windows
  - Only mechanism in OS, policy up to service
    - Low overhead mechanisms for simple services (window manager)
    - Weighty mechanisms for services (file system)
  - In memory logs are written to disk by transaction manager
    - Frequency is a performance consideration
    - Can be more opportunistic if failures are infrequent
12. Implementation Notes
- Log maintainence
    - Transaction manager write log records for recovering to persistent state
    - Frequency of “log force” impacts performance
  - Services have to choose mechanisms commensurate with their recovery requirements
    - Can impact all clients due to synchronous I/O requirements
13. Quicksilver Conclusion
- Ideas of transactions as a fundamental OS mechanism to bundle in state recovery of OS services found resurgence in the 90s with LRVM
    - In 2010, used to provide safeguard against system vulnerability and malicious attacks on a system in a research OS called Texas
  - Commercial OS are always focused on performance
    - Reliability takes a back seat
    - Writing to a file is only in memory until it's flushed to disk
  - Storage class memory
    - Nonvolatile, but with latency properties similar to DRAM