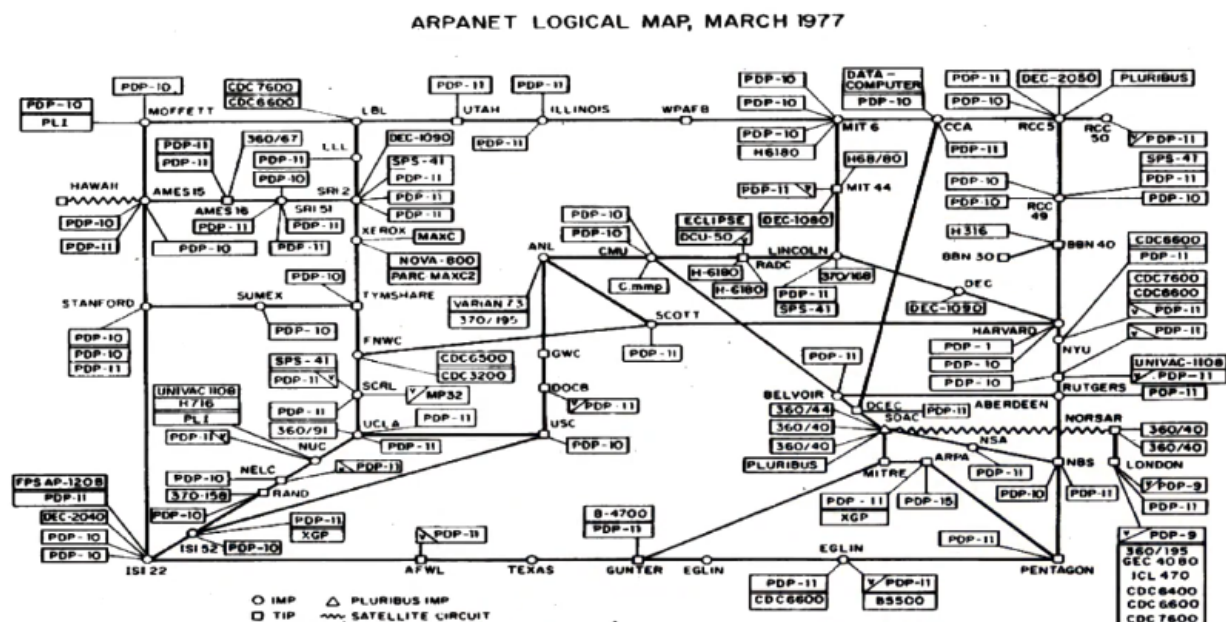


# Information Security

## Principles of Information Security

- \* Principles of Information Security Introduction
  - Introduction to terminologies pertaining to information security
  - Cryptographic techniques used to provide distributed authentication
    - + Private key infrastructure
- \* Firsts from Computing Pioneers
  - First vision for a network of computers: 1963
  - First computer-computer communication: 1969
  - First email: 1971
  - Main frame: Contained CPU, memory, and I/O with CRT terminals to connect to the mainframe in 1975 (timesharing system)



All of the Computers in the World (1977)


- \* Terminologies
  - When to release information?
    - + Privacy: When do individuals expect information to be protected or released?
    - + Security: Concerned with protection and authentication (system)
  - Comprehensive set of security concerns
    - + Unauthorized information release
    - + Unauthorized information modification
    - + Unauthorized denial of use (DOS attacks)
  - Goal of a secure system
    - + Prevent all violations of the concerns above
    - + Negative statement, impossible to achieve absolutely
    - + At best, negative statement provides false sense of security
- \* Levels of Protection
  - Unprotected
    - + MSDOS: Hooks for mistake prevention, not the same as security

- All or nothing
  - + IBM VM-370: Only way to interact with one another is explicit I/O across virtual machines
- Controlled sharing
  - + Access lists associated with files
- User programmed sharing controls
  - + Unix-like semantics for files (owner, group, everyone else)
- Strings on information
  - + "TOP SECRET" need to deal with dynamics of use
- \* Design Principles
  - Economy of mechanisms: Easy to verify whether it works or not
  - Fail safe defaults: Explicitly allow access -> default should not be no access (no way to guarantee information is protected)
  - Complete mediation: Security mechanism shouldn't take any shortcuts
    - + Caching passwords is a bad idea
  - Open design: Publish the design but protect the keys
    - + Must present keys to get information and breaking the keys should be computationally infeasible
    - + Detection is easier than prevention (don't know what attacks are possible, so impossible to prevent all of them)
  - Separation of privilege: Two keys to open a vault
  - Least privilege: "Need to know" based controls
    - + Require administrative privileges for certain things
    - + Origin for the idea of firewalls
  - Least common mechanism: Should mechanism be implemented in the kernel or as a library running on top?
  - Psychological acceptability: Good UI
  - Important takeaways
    1. All of the principles are positive statements
    2. All of the principles are still relevant to today's systems, despite being theorized at a time when computers weren't networked
    3. Difficult to crack protection boundary (computationally infeasible)
    4. Detection rather than prevention
- \* Principles of Information Security Conclusion
  - Paper by SALSA is a classic
  - Visionary ideas, thought of at a time when computers were independent

## Security in Andrew

- \* Security in Andrew Introduction
  - Enable students at CMU to walk up to any workstation and use it (1988)
  - Access files from central server
  - Assume network was untrusted
  - Focus on security using a private key infrastructure
- \* State of Computing Circa 1988
  - Local disks served as efficient caches
  - Vision of Andrew file system and Coda (both from CMU)
    - + Walk up to any workstation and log in
    - + Your content magically appears
    - + Similar to cloud and mobile devices today
- \* Andrew Architecture
  - Client workstation (virtues)
    - + Unix
    - + Connected to LAN through insecure network links

- + Servers within a "vice" can communicate securely
- Venus
  - + Process on each virtue for user authentication and client caching
  - + Users use RPC to transfer files from vice to virtue
  - + RPC must be secure for passing parameters and receiving results
- Client communication with servers must be encrypted
- Server communication with each other does not



Andrew Architecture

#### \* Encryption Primer

- Private key system
  - + Symmetric keys to encrypt and decrypt (e.g. passwords to login)
  - + Publish the design but protect the key (SALSA)
  - + Key distribution is difficult as the organization grows
  - + Sender:  $\text{data} \rightarrow \text{encrypt}(\text{data}, \text{key}) \rightarrow \text{cyphertext}$
  - + Receiver:  $\text{cyphertext} \rightarrow \text{decrypt}(\text{cyphertext}, \text{key}) \rightarrow \text{data}$
- Public key system
  - + Asymmetric keys (pair of keys)
  - + Public key published  $\rightarrow$  Encrypt
  - + Private key  $\rightarrow$  Decrypt
  - + Mathematical basis is a one-way function
  - + Sender:  $\text{Data} \rightarrow \text{encrypt}(\text{data}, \text{public key}) \rightarrow \text{cyphertext}$

- + Receiver: Cyphertext -> decrypt(data, private key) -> data
- \* Private Key Encryption System in Action
  - With private key encryption
    - + A and B have exchanged keys
    - + B needs to know identity of sender to know which key to use
    - + The identity of the sender is sent in cleartext
    - + KeyA can be the same as KeyB
- \* Challenges for Andrew System
  - Authenticate user
    - + How can you verify the identity of the person logging in?
  - Authenticate server
    - + How can you be sure the message is from the actual server?
  - Prevent replay attacks
    - + A person sniffing packets from the wire shouldn't be able to resend a packet and fool the sender or receiver?
  - Isolate users
    - + Community is shielded from the actions of other users
  - Andrew used a private key crypto system to protect RPC calls
    - + Key distribution is not as big of a challenge
    - + Identity of sender in cleartext
  - Traditional Unix: Username, password
    - + Overuse of usernames and passwords results in a security hole
    - + Violates the principle of protecting the keys
  - Dilemma: What to use as identity and private key?
- \* Andrew Solution
  - Username and password only for login
  - Use ephemeral ID and keys for subsequent Venus-vice communication
  - Three classes of client-server interaction
    1. Login: Username, password
    2. RPC session establishment: Open communication with file system
    3. File system access during session: Download files and work locally
  - 1 uses username and password
  - 2 and 3 uses ephemeral IDs and keys
- \* Login Process
  - User logs in with username and password to login server
    - + Authentication server is separate from authentication server
    - + Login process returns cleartoken and secrettoken
    - + This communication is secure
  - Cleartoken: Data structure
    - + Extract handshake key client (HKC)
    - + Use HKC as private key for establishing a new RPC session
  - Secrettoken: Cleartoken encrypted with key known only to vice
    - + Encryption is different from the HKC
    - + Unique for this login session (just a bitstream)
    - + Use as ephemeral client-id for this login session
    - + Only vice knows how to decrypt
  - Venus throws away clear and secret tokens at the end of the session
  - Bind mechanism is at the core of the Andrew file system
    - + Used for establishing secure communication
- \* RPC Session Establishment
  - After a user logs in, Venus establishes an RPC session on behalf of the client (bind client-server)
  - Client sends:
    - + clientID, E[Xr,HKC]

- + Xr is a random number for each RPC
- + HKC is extracted from the cleartoken
- Server sends:
  - +  $E[(Xr+1, Yr), HKS]$  (HKC = HKS by design)
  - + Xr+1 establishes that the server is genuine
  - + Xr is encrypted, so this authenticates the server and prevents replay attacks
  - + Yr is another random number generated by the server
- Client sends:
  - +  $E[Yr+1, HKC]$
  - + Server checks if the value == Yr+1 to establish that the client is genuine and prevents replay attacks
- Server only uses HKC for establishing a login session
  - + Within an RPC session, a client may make many calls (open, read, write, close, etc)
  - + After the server validates the ID of the client, it creates a session key (SK) and sends it to the client
- Server sends:
  - +  $E[(SK, num), HKS]$
  - + Client can extract the SK using its own HKC and use it for the duration of the session
  - + num is the starting sequence number for RPC session (safeguard against replay attacks)
  - + Use SK as handshake key for the rest of RPC session with server
- \* Sequence Establishment
  - The sequence client(Xr) -> server(Xr+1) -> client establishes that the server is genuine
  - The sequence server(Yr) -> client(Yr+1) -> server establishes that the client is genuine
- \* Login is a Special Case of Bind
  - Client/server validation is identical for login and bind
  - Password used as HKC
  - Username used for clientID
  - Get back two tokens
    - + Secrettoken (encrypted with password)
    - + Cleartoken (encrypted with password)
  - Tokens kept by Venus for this login session
- \* Putting it All Together
  - Login using username and password
  - Vice sends secret and clear tokens (1)
  - Venus establishes an RPC session on behalf of the client using secret token and HKC
  - Vice sends session key for this particular RPC session (2)
  - Venus uses secret token and session key to make file system calls (3)
  - Upshots:
    - + username, password are only exposed once per login session
    - + HKC used only for new RPC session establishment
    - + SK used for all RPC calls to file system (valid for duration of RPC session)
- \* AFS Security Report Card
  - Mutual suspicion: Yes
    - Protection from fellow users
  - Protection from system for users: No
    - + Users must trust system, so no protection

- Confinement of resource usage: No
  - + User can make many calls on server and consume bandwidth (DoS)
- Authentication: Yes
  - + Validation of client-server
- Server integrity: No
  - + Servers are within a firewall, so if somebody penetrated this, they could wreak havoc
  - + Must physically and socially protect servers
- \* Security in Andrew Conclusion
  - Main takeaway: How do OS designers make the best decisions for information security to make a secure and usable distributed system?
    - + Able to compare against the principles outlined by SALSA
  - Introduces notion of audit chain for system administrators modifying the system
  - Access lists for files with positive and negative rights