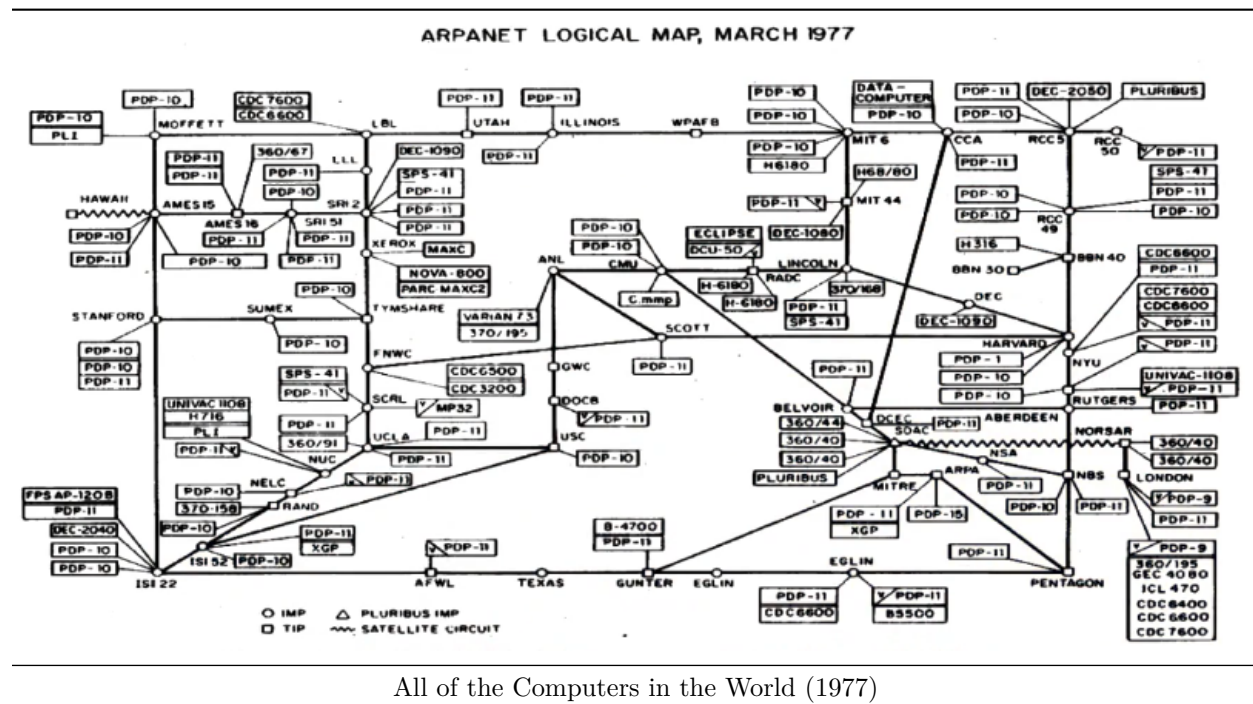


# Information Security

## Principles of Information Security

1. Principles of Information Security Introduction
  - Introduction to terminologies pertaining to information security
  - Cryptographic techniques used to provide distributed authentication
    - Private key infrastructure
2. Firsts from Computing Pioneers
  - First vision for a network of computers: 1963
  - First computer-computer communication: 1969
  - First email: 1971
  - Main frame: Contained CPU, memory, and I/O with CRT terminals to connect to the mainframe in 1975 (timesharing system)




3. Terminologies
  - When to release information?
    - Privacy: When do individuals expect information to be protected or released?
    - Security: Concerned with protection and authentication (system)
  - Comprehensive set of security concerns
    - Unauthorized information release
    - Unauthorized information modification
    - Unauthorized denial of use (DOS attacks)
  - Goal of a secure system
    - Prevent all violations of the concerns above
    - Negative statement, impossible to achieve absolutely
    - At best, negative statement provides false sense of security
4. Levels of Protection
  - Unprotected
    - MSDOS: Hooks for mistake prevention, not the same as security
  - All or nothing
    - IBM VM-370: Only way to interact with one another is explicit I/O across virtual machines

- Controlled sharing
    - Access lists associated with files
  - User programmed sharing controls
    - Unix-like semantics for files (owner, group, everyone else)
  - Strings on information
    - “TOP SECRET” need to deal with dynamics of use
5. Design Principles
- Economy of mechanisms: Easy to verify whether it works or not
  - Fail safe defaults: Explicitly allow access -> default should not be no access (no way to guarantee information is protected)
  - Complete mediation: Security mechanism shouldn’t take any shortcuts
    - Caching passwords is a bad idea
  - Open design: Publish the design but protect the keys
    - Must present keys to get information and breaking the keys should be computationally infeasible
    - Detection is easier than prevention (don’t know what attacks are possible, so impossible to prevent all of them)
  - Separation of privilege: Two keys to open a vault
  - Least privilege: “Need to know” based controls
    - Require administrative privileges for certain things
    - Origin for the idea of firewalls
  - Least common mechanism: Should mechanism be implemented in the kernel or as a library running on top?
  - Psychological acceptability: Good UI
  - Important takeaways
    - All of the principles are positive statements
    - All of the principles are still relevant to today’s systems, despite being theorized at a time when computers weren’t networked
    - Difficult to crack protection boundary (computationally infeasible)
    - Detection rather than prevention
6. Principles of Information Security Conclusion
- Paper by SALSA is a classic
  - Visionary ideas, thought of at a time when computers were independent

## Security in Andrew

1. Security in Andrew Introduction
  - Enable students at CMU to walk up to any workstation and use it (1988)
  - Access files from central server
  - Assume network was untrusted
  - Focus on security using a private key infrastructure
2. State of Computing Circa 1988
  - Local disks served as efficient caches
  - Vision of Andrew file system and Coda (both from CMU)
    - Walk up to any workstation and log in
    - Your content magically appears
    - Similar to cloud and mobile devices today
3. Andrew Architecture
  - Client workstation (virtues)
    - Unix
    - Connected to LAN through insecure network links
    - Servers within a “vice” can communicate securely
  - Venus
    - Process on each virtue for user authentication and client caching

- Users use RPC to transfer files from vice to virtue
- RPC must be secure for passing parameters and receiving results
- Client communication with servers must be encrypted
- Server communication with each other does not



Andrew Architecture

#### 4. Encryption Primer

- Private key system
  - Symmetric keys to encrypt and decrypt (e.g. passwords to login)
  - Publish the design but protect the key (SALSA)
  - Key distribution is difficult as the organization grows
  - Sender:  $\text{data} \rightarrow \text{encrypt}(\text{data}, \text{key}) \rightarrow \text{cyphertext}$
  - Receiver:  $\text{cyphertext} \rightarrow \text{decrypt}(\text{cyphertext}, \text{key}) \rightarrow \text{data}$
- Public key system
  - Asymmetric keys (pair of keys)
  - Public key published  $\rightarrow$  Encrypt
  - Private key  $\rightarrow$  Decrypt
  - Mathematical basis is a one-way function
  - Sender:  $\text{Data} \rightarrow \text{encrypt}(\text{data}, \text{public key}) \rightarrow \text{cyphertext}$
  - Receiver:  $\text{Cyphertext} \rightarrow \text{decrypt}(\text{data}, \text{private key}) \rightarrow \text{data}$

#### 5. Private Key Encryption System in Action

- With private key encryption

- A and B have exchanged keys
  - B needs to know identity of sender to know which key to use
  - The identity of the sender is sent in cleartext
  - KeyA can be the same as KeyB
6. Challenges for Andrew System
- Authenticate user
    - How can you verify the identity of the person logging in?
  - Authenticate server
    - How can you be sure the message is from the actual server?
  - Prevent replay attacks
    - A person sniffing packets from the wire shouldn't be able to resend a packet and fool the sender or receiver?
  - Isolate users
    - Community is shielded from the actions of other users
  - Andrew used a private key crypto system to protect RPC calls
    - Key distribution is not as big of a challenge
    - Identity of sender in cleartext
  - Traditional Unix: Username, password
    - Overuse of usernames and passwords results in a security hole
    - Violates the principle of protecting the keys
  - Dilemma: What to use as identity and private key?
7. Andrew Solution
- Username and password only for login
  - Use ephemeral ID and keys for subsequent Venus-vice communication
  - Three classes of client-server interaction
    - Login: Username, password
    - RPC session establishment: Open communication with file system
    - File system access during session: Download files and work locally
  - 1 uses username and password
  - 2 and 3 uses ephemeral IDs and keys
8. Login Process
- User logs in with username and password to login server
    - Authentication server is separate from authentication server
    - Login process returns cleartoken and secrettoken
    - This communication is secure
  - Cleartoken: Data structure
    - Extract handshake key client (HKC)
    - Use HKC as private key for establishing a new RPC session
  - Secrettoken: Cleartoken encrypted with key known only to vice
    - Encryption is different from the HKC
    - Unique for this login session (just a bitstream)
    - Use as ephemeral client-id for this login session
    - Only vice knows how to decrypt
  - Venus throws away clear and secret tokens at the end of the session
  - Bind mechanism is at the core of the Andrew file system
    - Used for establishing secure communication
9. RPC Session Establishment
- After a user logs in, Venus establishes an RPC session on behalf of the client (bind client-server)
  - Client sends:
    - clientID,  $E[X_r, \text{HKC}]$
    - $X_r$  is a random number for each RPC
    - HKC is extracted from the cleartoken
  - Server sends:
    - $E[(X_r+1, Y_r), \text{HKS}]$  (HKC = HKS by design)

- Xr+1 establishes that the server is genuine
    - Xr is encrypted, so this authenticates the server and prevents replay attacks
    - Yr is another random number generated by the server
  - Client sends:
    - E[Yr+1,HKC]
    - Server checks if the value == Yr+1 to establish that the client is genuine and prevents replay attacks
  - Server only uses HKC for establishing a login session
    - Within an RPC session, a client may make many calls (open, read, write, close, etc)
    - After the server validates the ID of the client, it creates a session key (SK) and sends it to the client
  - Server sends:
    - E[(SK,num),HKS]
    - Client can extract the SK using its own HKC and use it for the duration of the session
    - num is the starting sequence number for RPC session (safeguard against replay attacks)
    - Use SK as handshake key for the rest of RPC session with server
10. Sequence Establishment
- The sequence client(Xr) -> server(Xr+1) -> client establishes that the server is genuine
  - The sequence server(Yr) -> client(Yr+1) -> server establishes that the client is genuine
11. Login is a Special Case of Bind
- Client/server validation is identical for login and bind
  - Password used as HKC
  - Username used for clientID
  - Get back two tokens
    - Secrettoken (encrypted with password)
    - Cleartoken (encrypted with password)
  - Tokens kept by Venus for this login session
12. Putting it All Together
- Login using username and password
  - Vice sends secret and clear tokens (1)
  - Venus establishes an RPC session on behalf of the client using secret token and HKC
  - Vice sends session key for this particular RPC session (2)
  - Venus uses secret token and session key to make file system calls (3)
  - Upshots:
    - username, password are only exposed once per login session
    - HKC used only for new RPC session establishment
    - SK used for all RPC calls to file system (valid for duration of RPC session)
13. AFS Security Report Card
- Mutual suspicion: Yes
    - Protection from fellow users
  - Protection from system for users: No
    - Users must trust system, so no protection
  - Confinement of resource usage: No
    - User can make many calls on server and consume bandwidth (DoS)
  - Authentication: Yes
    - Validation of client-server
  - Server integrity: No
    - Servers are within a firewall, so if somebody penetrated this, they could wreak havoc
    - Must physically and socially protect servers
14. Security in Andrew Conclusion
- Main takeaway: How do OS designers make the best decisions for information security to make a secure and usable distributed system?
    - Able to compare against the principles outlined by SALSA
  - Introduces notion of audit chain for system administrators modifying the system

- Access lists for files with positive and negative rights