

Realtime and Multimedia

Time Sensitive-Linux

1. Time Sensitive-Linux Introduction

- Traditional operating systems catered to databases and scientific applications which are throughput oriented
- Modern workloads support A/V and video games, which have more strict latency requirements and real time guarantees
- How do you provide guarantees for real time applications in the presence of background throughput-oriented applications such as databases?
- How do you bound the performance loss of throughput-oriented applications in the presence of latency-sensitive applications?

2. Sources of Latency

- Time sensitive applications require quickly responding to an event
 - Playing a video game
- Three sources of latency for time-sensitive events
 - Timer latency: Inaccuracy of timing mechanism; delta between timer event and timer interrupt due to the granularity of the timing mechanism (10ms in Linux)
 - Preemption latency: Interrupt could occur when kernel is in the middle of doing something and cannot be preempted (handler disabled)
 - Scheduler latency: Already a higher priority task in the scheduling queue
- Latency = Activation of event - Happening of event
 - Latency = $T_a - T_h$

3. Timers Available

| Choice of Timer | Pro | Con |
|-----------------|------------------|---------------------------|
| Periodic | Periodicity | Event recognition latency |
| One-shot | Timely | Overhead |
| Soft | Reduced overhead | Polling overhead, latency |
| Firm | Combines all | |


4. Firm Timer Design

- Goal: accurate timing with low overhead
- Combine one-shot and soft timers
- Overshoot: Knob between hard and soft timers
 - Time between timer expiring and interrupt being triggered
 - If a syscall occurs during this period, we dispatch the expired timers and reprogram the one-shot timers
 - Advantage: The point at which the timer is programmed to interrupt will not actually cause an interrupt
- Firm timer provides accuracy without additional overhead

5. Firm Timer Implementation

- Task: Schedulable entity (T_1, T_2, T_3)
- Timer-q data structure: Queue of tasks sorted by expiry time
- APIC timer hardware
 - APIC: Advanced Programmable Interrupt Controller
 - Reprogramming a one-shot timer only takes a few cycles
 - Sets a value in a register and decrements it each cycle
 - Allows for the implementation of very fine-grained timers
- When the APIC timer expires, the interrupt handler traverses timer-q data structure and look for tasks whose timers have expired
 - Callback handler associated with each event are called
 - Tasks are removed from the timer-q

- If the task is periodic, the handler will re-enqueue the task
 - Soft timers: Eliminates need for fielding one-shot interrupts
 - If there's a long time between one-shot events and many periodic events occurring in between, we can process an upcoming one-shot event at the preceding periodic event
 - Processing an event early is fine, just don't want to be late
 - Periodic events are much more efficient in the kernel ($O(1)$)
 - One-shot events are $O(\log N)$ where N is the number of timers
 - Reduces the timer latency of event handling
6. Reducing Kernel Preemption Latency
- Approaches
 - Explicit insertion of preemption points in kernel (look for events and take action)
 - Allow preemption anytime the kernel is not manipulating shared data structures (prevents race conditions)
 - Lock-Breaking Preemptive Kernel
 - Combines the above ideas to reduce spin lock holding time
 - Replace long critical section with two shorter critical sections
 - $\text{acq}() \rightarrow \text{rel}()$ vs $\text{acq}() \rightarrow \text{rel}() \rightarrow \text{reacq}() \rightarrow \text{rel}()$
 - Can preempt the kernel after the first release
7. Reducing Scheduling Latency
- Proportional period scheduler
 - Can reserve a portion of the time quantum for throughput-oriented tasks to guarantee that their requirements are still satisfied in a timely manner
 - T is a time window/quantum (exposed to an application)
 - Task requests Q , a proportion of T
 - Both are adjustable using a feedback control mechanism
 - Scheduler attempts to satisfy multiple requests within a quantum
 - If one task needs $2/3 T$ and another needs $1/3 T$, both can be satisfied within a single quantum
 - Priority inversion
 - High priority $C1$ task issues a blocking call to a low priority server
 - Until the server finishes, high priority task cannot proceed
 - While the low priority task is executing, another task $C2$ can become runnable which will preempt the lower priority (priority $C1 > C2$)
 - This results in priority inversion for $C1$
 - Priority-based Scheduling
 - Server assumes the priority of $C1$ while it is executing its task
 - Therefore, no priority inversion occurs
 - How TS-Linux deals with time-sensitive tasks
 - Firm timer design that increases the accuracy of the timer without additional overhead
 - Using a preemptable kernel to reduce kernel latency
 - Using priority based scheduling to avoid priority and inversion and guaranteeing a portion of the CPU for throughput-oriented tasks



Proportional Period Scheduler

8. Time Sensitive-Linux Conclusion

- TS-Linux is able to provide quality of service guarantees for realtime applications running on commodity operating systems
- Admission control using proportional period scheduling ensures that throughput-oriented tasks receive adequate time to execute

Persistent Temporal Streams

1. Persistent Temporal Streams Introduction

- Focus on middleware between commodity operating systems and realtime, distributed multimedia applications


2. Programming Paradigms

- Parallel programs: Pthreads
- Distributed programs: Sockets (NFS, RPC)

3. Novel Multimedia Apps

- Novel distributed applications tend to be sensor-based
 - Cameras, radars, microphones, temperature sensors, etc.
 - Access sensor data through the Internet
 - Want to do live stream analysis of sensor data
 - Sense -> Prioritize -> Process -> Activate actuators

- These computations are intensive and real-time
 - Want to reduce latency as much as possible
4. Example - Large Scale Situational Awareness
- Monitoring activities in an airport, anything anomalous causes an alert
 - London has over 400,000 cameras to monitor
 - Need to prune the data to deal with this many sensors
 - Can't rely on humans to monitor thousands of cameras
 - Avoid false positives and false negatives
 - Programming infrastructure must facilitate domain expert
 - Pain points
 - Right level of abstraction to promote ease of use
 - Propagate temporal causality
 - Correlate live data with historical data to draw inferences
5. Programming Model for Situational Awareness
- Sequential program for video analysis
 - Camera -> Detection -> Tracking -> Recognition -> Alarm
 - Objective in situational aware applications: Process streams of data for high level inference (not watching a video)
 - How does this scale to thousands of cameras?
 - PTS is an example of a distributed programming system catering to the needs of situational awareness applications
6. PTS Programming Model
- Similar to Unix socket API
 - Distributed application with threads and channels
 - Channel holds time-sequenced data objects, allows many-to-many connections
 - Thread is generating time-sequenced data objects and inserting into the channel
 - Can be multiple producers inserting into a channel and multiple consumers receiving from a channel
 - Channel shows temporal evolution of what is being produced by a particular thread (one frame from a camera)
 - `put(item,)`
 - `get(<lowerbound, upperbound>)`
 - Can specify time in an abstract way (oldest, newest) in addition to explicit lower/upper bounds
 - `Channel ch1 = lookup("VideoChannel");`
 - `while(true) {`
 - `response r = ch1.get(<LB, UB>);`
 - `// process data`
 - `// produce output`
 - `ch2.put(item,); }`
 - PTS model allows for the ability to associate timestamp with data items produced by a computation
 - Allows for temporal causality
 - Every stream is temporally indexed, so a computation can correlate incoming streams and recognize which data items are temporally related to one another



Proportional Period Scheduler


7. Bundling Streams

- Might need data from multiple data streams simultaneously
- Can create a “stream group” which is a bundle of different streams
 - Video, audio, text, gesture
 - One is the anchor stream, the other are dependents
- groupget: Get corresponding time-stamped items from all the streams in the group
 - Gets temporally related items from all streams in a group at once


8. Power of Simplicity

- Interpose channels between modules to get/put data
- Convert from sequential program to distributed using the channel abstraction and get/put primitives

Sequential program for video analytics



Distributed program with get/put between modules

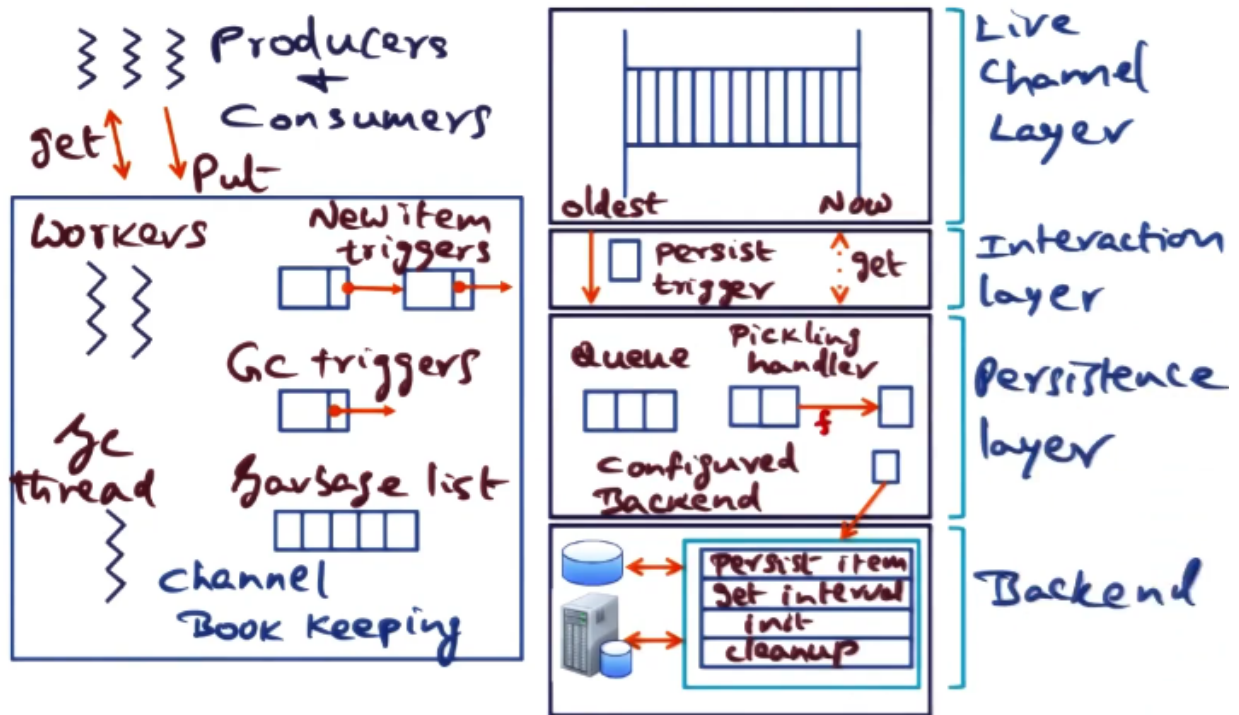


PTS Program Workflow

9. PTS Design Principles

- PTS provides simple abstractions and interfaces
 - Channel and get/put

- Do the heavy lifting (systems) under the covers
 - PTS channels
 - Can be anywhere (similar to Unix sockets)
 - Can be accessed from anywhere (similar to Unix sockets)
 - Network-wide unique (similar to Unix sockets)
 - Time is a first class entity: Time is manipulated by the application in the way it specifies items to the runtime system and queries the runtime system using time as the index into the channel
 - Persistent streams under application control: Data is continuously being produced, can't hold it all in memory. Stored to archival storage
 - Seamlessly handle live and historical data
 - PTS is simple to use, but provides the necessary handles to make the life of a domain expert developing a situation awareness application
10. Persistent Channel Architecture
- How does PTS support this simple get/put programming model?
 - Everything is either a producer or a consumer
 - Threads
 - Workers: Produce new data
 - Garbage collection: Delete data from garbage list
 - Triggers
 - New item triggers occur every time a producer in a worker thread creates a new piece of data
 - Garbage collection triggers: Might specify channel to only keep last 30 seconds of data, throw everything else away. Move data into garbage list (no longer relevant to application)
 - Persist triggers indicate that an item has become old and should be moved from memory to archive
 - Implementation of channel is a three-layer architecture
 - Top: Live channel layer; reacts to new item triggers coming from worker threads
 - Middle: Interaction layer; Go between the live channel and persistence layer
 - Bottom: Persistence layer; receives persistence triggers and moves items from the channel according to the pickling handler. Persisted items must go to some non-volatile storage device
 - Pickling handler: Developer specifies a function to be used every time some of the items are persisted (might condense in some way)
 - Backend layer determines how to keep data in non-volatile storage
 - Application choice as to which backend layer it wants to use
 - Backend layers: MySQL, Unix filesystem, IBM GPFS



PTS Persistent Channel Architecture

11. PTS Conclusion

- Map reduce is to big data applications as PTS is to live stream analysis applications
- Unique features
 - Time-based distributed data structures for streams
 - Automatic data management
 - Transparent stream persistence