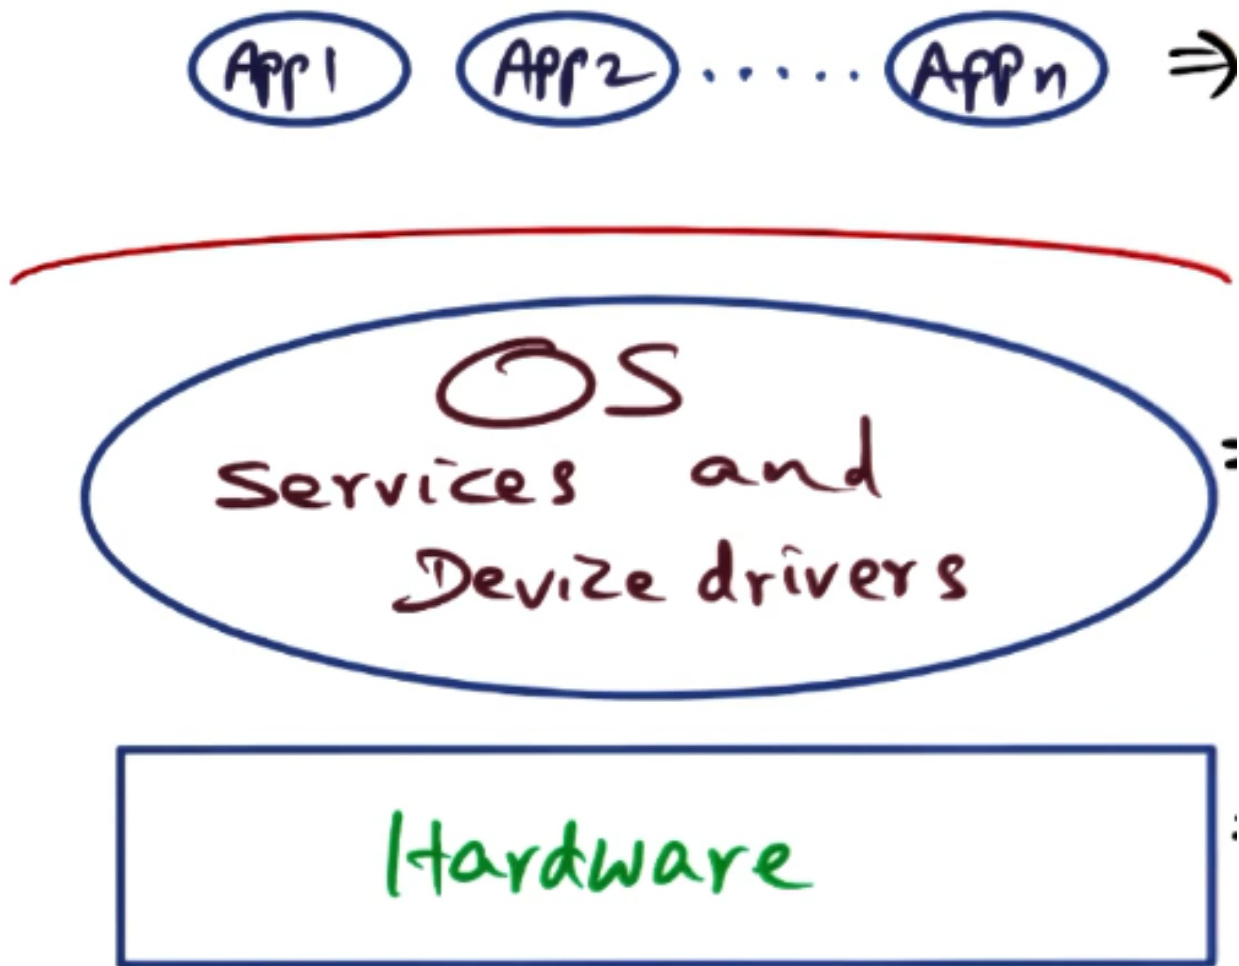


Operating System Structure

OS Structure Overview

1. Operating System Services
 - Process/thread management and scheduling
 - Memory management (protection, sharing, demand paging)
 - Interprocess communication (IPC)
 - File system
 - Access to I/O devices such as microphones and speakers
 - Access to the network
2. OS Structure: The way the operating system software is organized with respect to the applications that it serves and the underlying hardware that it manages
3. Goals of OS Structure
 - Protection: Within and across users and to the OS itself
 - Performance: Time taken to perform the services
 - Flexibility: Extensibility (not one size fits all)
 - Scalability: Performance improves with more hardware resources
 - Agility: Adapting to changes in application needs and/or resource availability
 - Responsiveness: Reacting to external events
4. Commercial OS: Don't meet all goals laid out above
 - Linux
 - MacOS
 - Windows
5. Monolithic Structure
 - Each app has its own hardware address space
 - OS also has its own hardware address space (protected from apps)
 - Hardware is managed by the OS
 - OS services and device drivers are all contained in one blob
 - Reduces performance loss by consolidation
 - All of the components are aware of each other and interactions can be optimized
 - Con: Less extensible (no customization, useful for different apps)
 - Video games: Responsiveness is desired
 - Computing prime numbers: Sustained CPU time is desired



Monolithic OS Structure

6. DOS-like Structure

- Disk operating system
- No protection layer between the application and the OS
- Pro: Improved performance (access to system services is like a procedure call)
- Con: Decreased protection (an errant application can corrupt the OS)
- In the early days of computing, only one process could run at a time
 - Simplicity and performance took precedence
 - No overhead for system calls vs developer procedures
- Loss of protection is unacceptable for a general-purpose OS


7. Opportunities for Customization

- Memory management (page fault)
 - OS finds a free page frame (page replacement algorithm)
 - OS updates page table for faulty process
 - Resume process
- The page replacement algorithm can't be ideal for every application
- Can customize how page replacement is handled

8. Microkernel Structure

- Each app in its own hardware address space
- Microkernel runs in privileged mode and provides simple mechanisms

- Threads
- Address space
- IPC
- OS services are implemented on top of microkernel (same privilege as applications)
- No distinction between OS services and applications in terms of how they run
- Pro: Extensibility; services can be easily replaced
- Con: Performance; in a monolithic structure, filesystem and memory management are all within the kernel. In a microkernel structure, an app might go through the kernel, to the filesystem service, back to the microkernel, and back to the app. All of this is IPC overhead.
- Border crossings
 - User space \leftrightarrow system space copying
 - Explicit costs in changing context
 - Implicit costs due to cache misses (change in locality)

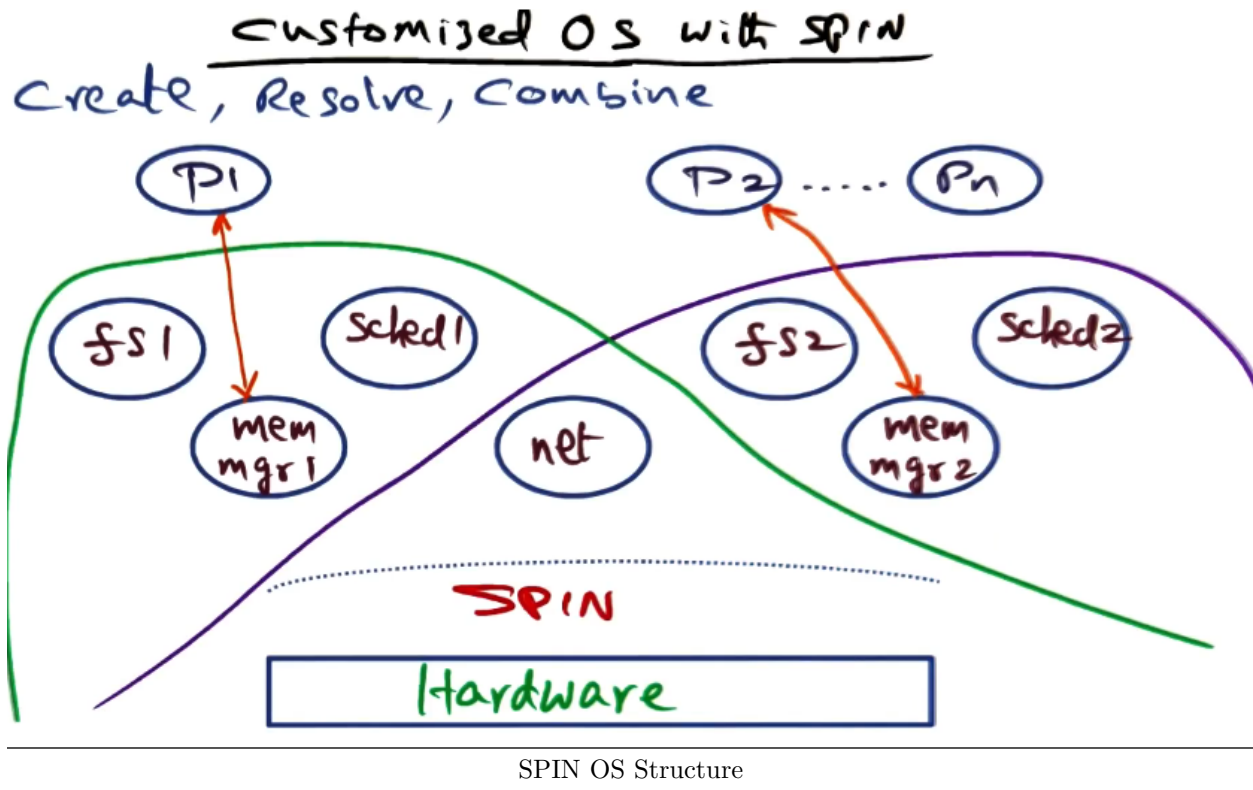


Microkernel OS Structure

Feature	Monolithic	DOS-like	Microkernel
Extensibility		X	X
Protection	X		X
Performance	X	X	

The SPIN Approach

1. Two Premises:
 - Microkernel design compromises on performance due to frequent border crossings
 - Monolithic design does not lend itself to extensibility
2. Goals
 - Thin (like microkernel) -> only mechanisms, not policies
 - Access to resources without border crossing (like DOS)
 - Flexibility for resource management (like microkernel) without sacrificing protection and performance (like monolithic)
3. Approaches to Extensibility
 - Hydra OS
 - Kernel mechanisms for resource allocation
 - Capability-based resource access
 - Resource management as coarse-grained objects to reduce border crossing overhead
 - Mach OS
 - Microkernel-based
 - Focused on extensibility and portability
 - Performance wasn't prioritized
4. SPIN's Approach to Extensibility
 - Co-location of kernel and extensions to avoid border crossings
 - Compiler enforced modularity (strongly typed language, Modula)
 - Logical protection domains (not hardware address spaces)
 - Dynamic call binding -> flexibility
 - Make extensions as cheap as a procedure call
5. Logical Protection Domains
 - Modula-3 safety and encapsulation mechanisms
 - Type safety, automatic storage management
 - Objects, threads, exceptions, generic interfaces
 - Fine-grained protection via capabilities
 - Hardware resources (page frame)
 - Interfaces (page allocation module)
 - Collection of interfaces (entire virtual memory subsystem)
 - Capabilities implemented as language supported pointers
 - Modula-3 pointers are type-specific (no casting)
6. SPIN Mechanisms for Protection Domains
 - Create: initialize with object file contents and export names that are contained as entry point methods into the object to be externally visible
 - Resolve: Names between source and target domains
 - Once resolved, resource sharing at memory speeds
 - Combine: To create an aggregate domain from smaller domains
7. Customized OS with SPIN
 - No border crossing between services and mechanisms provided by SPIN

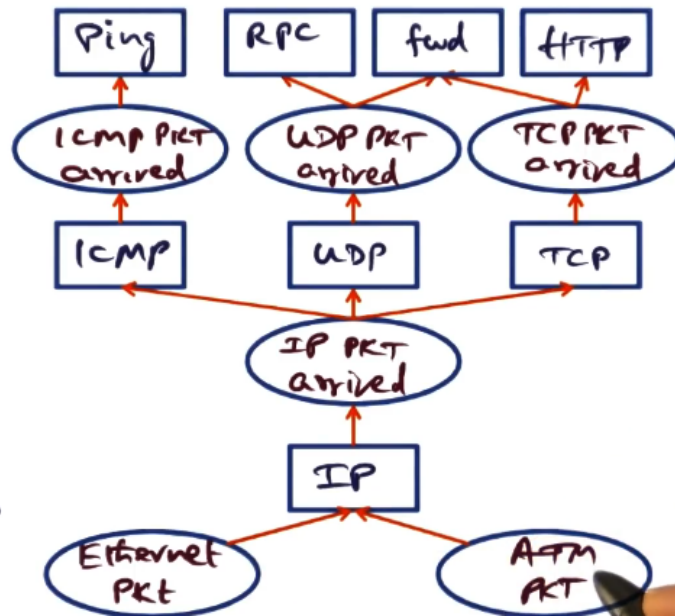



8. SPIN Mechanisms for Events

- OS needs to be able to field events (interrupts, system calls)
- Event-based communication model
 - Services register event handlers (one-to-one, one-to-many, many-to-one)

SPIN Mechanisms for Events

Event-based Communication model



SPIN Event Handlers

9. Default Core Services in SPIN

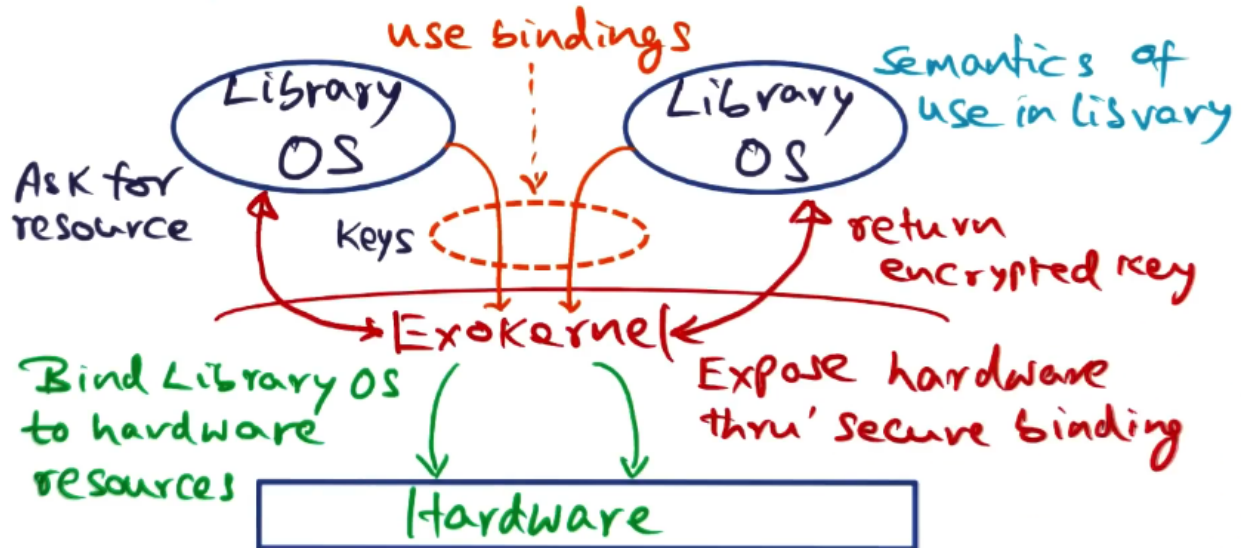
- SPIN provides interface procedures for implementing these services
- Automatically invoked when the hardware event occurs
- Memory management
 - Physical address (allocate, deallocate, reclaim)
 - Virtual address (allocate, deallocate)
 - Translation (create/destroy address space, add/remove mapping)
 - Event handlers (page fault, access fault, bad address)
- CPU Scheduling
 - SPIN abstraction: strand (semantics defined by extension)
 - Event handlers (block, unblock, checkpoint, resume)
 - SPIN global scheduler (interacts with application threads package)
- Core services are trusted because they provide access to hardware mechanisms
 - Services may need to step outside the language-enforced protection model to control the resources
 - Applications that run on top of extensions must trust extension
 - However, these are isolated (don't impact things that don't rely on the extension)

The Exokernel Approach

1. Exokernel: Kernel exposes hardware explicitly to OS extensions
 - Decouple authorization of hardware from actual use
 - Expose hardware through secure binding (binds library OS to hardware resources)
 - Gives encrypted key to library OS; semantics of use are up to library
 - Can't be forged or transferred
 - Library gives key to exokernel to validate it can use the resource

Exokernel Approach to Extensibility

Decouple Authorization from use



Exokernel OS Structure

2. Examples of Candidate Resources

- TLB Entry
 - Virtual to physical mapping done by library
 - Binding presented to exokernel
 - Exokernel puts it into hardware TLB
 - Process in library OS uses multiple time without exokernel intervention
- Packet filter
 - Predicates loaded into kernel by library OS
 - Check on packet arrival by Exokernel

3. Implementing Secure Bindings

- Hardware mechanisms (TLB entry, physical page frame, frame buffer)
- Software caching
 - “Shadow” TLB in software for each library OS
 - Exokernel dumps hardware TLB into software TLB data structure on context switch
- Downloading code into kernel
 - Functionally equivalent to SPIN extensions
 - Compromises protection more than SPIN; as long as SPIN’s logical protection domains are following Modula-3’s compile and runtime verification, no compromise. Exokernel doesn’t make this guarantee

4. Memory Management in Exokernel

- When a thread incurs a page fault, it’s caught by Exokernel
- Then, it’s passed to the library OS through the registered handler
- Library services fault (might require requesting a page frame from Exokernel)


5. Memory Management Using Software TLB

- On context switch, much of performance hit comes from loss of cache locality
- To mitigate this overhead, Exokernel implements a software TLB
- Software TLB: Data structure representing a snapshot of the hardware TLB; captured on context switch for the library OS

- On context switch, Exokernel preloads TLB with software TLB

Default Core Services in Exokernel

Memory management



Exokernel Memory Management


6. CPU Scheduling in Exokernel

- Linear vector of "time slots"
- Each library OS will get one time quantum at a time
- When timing interrupt occurs, Exokernel will switch to next library OS
- Some time between switching when library OS gets to save state
- Exokernel only interferes for page faults

Default Core Services in Exokernel

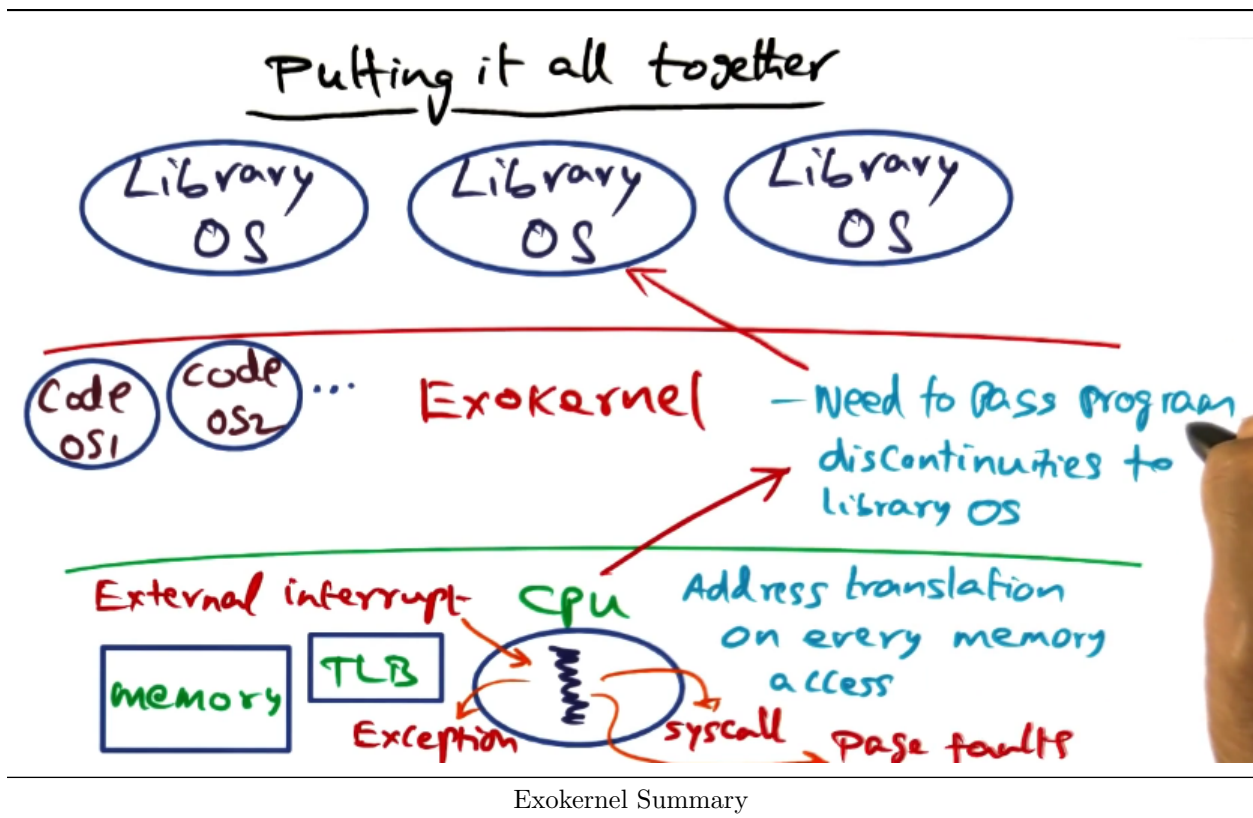
CPU Scheduling

— Linear vector of "time slots"



Exokernel CPU Scheduling

7. Secure Binding - How can a library OS securely insert code into the kernel?
 - Performance optimization to avoid border crossings
 - In both SPIN and Exokernel, the ability to add extensions must be restricted to a trusted set of users
8. Revocation of Resources
 - Exokernel needs a way of revoking resources (space/time) that have been allocated to a library OS
 - Revoke call is an upcall into library OS (repossession vector)
 - Library OS must clean up when it receives a repossession vector
 - Stash data in a page to disk when Exokernel reclaims
 - Library can "seed" Exokernel for "autosave"
9. Code Usage by Exokernel - Library OS gives code to run on specific events
 - Packet filter for de-multiplexing of network packets
 - Run code in Exokernel on behalf of library OS not currently scheduled (e.g., garbage collection for an application)
10. Summary; Achieving extensibility, protection, and performance
 - Performance: Exokernel allows for performance critical code in library OS to be downloaded securely into Exokernel
 - Extensibility: Exokernel exposes hardware resources to library OS to allow them to implement functionality however is needed
 - Protection: Exokernel provides encrypted keys to ensure that only the intended library OS can access hardware
 - Exokernel must be involved when interrupts occur



11. Exokernel Data Structures

- PE data structure: Handler entry points for different event types
 - Exceptions, interrupts, system calls, memory mappings
- Software TLB to preserve page mappings
 - Must specify a set of guaranteed mappings for Exokernel to maintain
 - Not all hardware TLB entries, just the specified ones
- Similar to event handler structure in SPIN OS

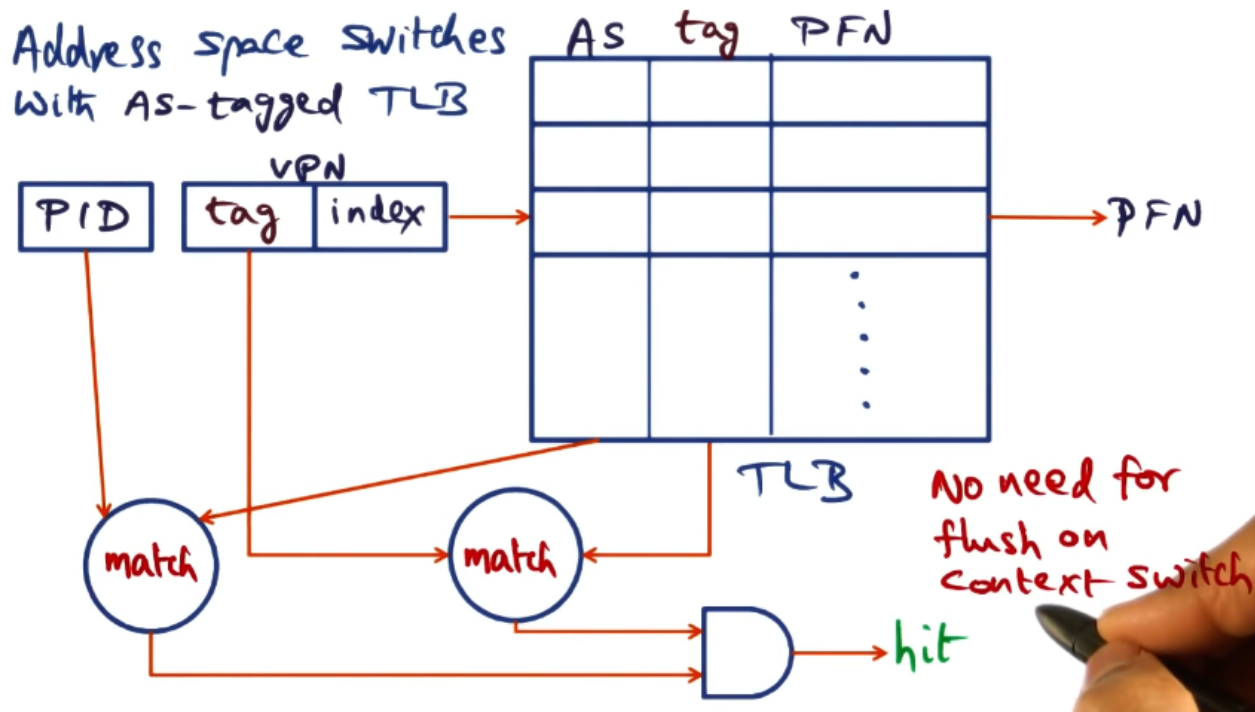
12. Performance Results of SPIN and Exokernel

- Can qualitatively argue that extensibility is provided, but must prove quantitatively that performance isn't being sacrificed
- UNIX: Monolithic
- Mach: Microkernel (CMU)
- Both SPIN and Exokernel outperform Mach regarding changing protection domains
- Both SPIN and Exokernel do as well as Unix for dealing with system calls

The L3 Microkernel Approach

1. SPIN and Exokernel were comparing against CMU's Mach, a microkernel approach with a focus on portability
 - However, can we design a microkernel with a focus on performance?
2. Microkernel: Each service in its own address space
 - Applications sit on top of OS services
 - OS services: Memory management, file system, storage module
 - Microkernel provides simple abstractions (address space, IPC)
3. Potentials for Performance Loss
 - Border crossings: Implicit and explicit costs
 - Applications and microkernel in different privilege levels
 - System services have to consult other services (cross address space)

- Explicit cost « implicit cost
- TLB flush takes 864 cycles in Pentium processor
- Small protection domains
 - Switches can be made efficient by careful construction
- Large protection domains
 - Switching cost not important
 - Cache effects and TLB effects dominate



Address Space Tagged TLB

- Thread Switches and IPC
 - Shown to be competitive to SPIN and Exokernel by construction
 - Switch involves saving all volatile state of processor
- Memory Effects
 - Assumption is that loss of locality in a microkernel-based design » monolithic design
 - Liedtke says that splitting the hardware address space into small protection domains using segment registers will result in warmer caches on context switches
 - Each process only occupies a small memory footprint, which results in a smaller footprint in the caches as well
 - Unavoidable if protection domains are large, regardless of how the operating system is designed
- Mach's Expensive Border Crossing
 - Focus on portability (can run on different processor architectures)
 - Code bloat -> large memory footprint -> Less locality -> more cache misses -> longer latency for border crossing
 - Kernel memory footprint is the culprit, not microkernel itself
- Thesis of L3 for OS Structuring
 - Minimal abstractions in microkernel
 - Address spaces, threads, IPC, unique IDs
 - These four abstractions are needed by any subsystem that provides functionality to end users
 - Microkernels are processor-specific in implementation (not portable)

- Right set of microkernel abstractions and processor-specific implementation -> efficient processor-independent abstractions at higher layers
 - Processor-dependent kernel with processor-independent abstractions