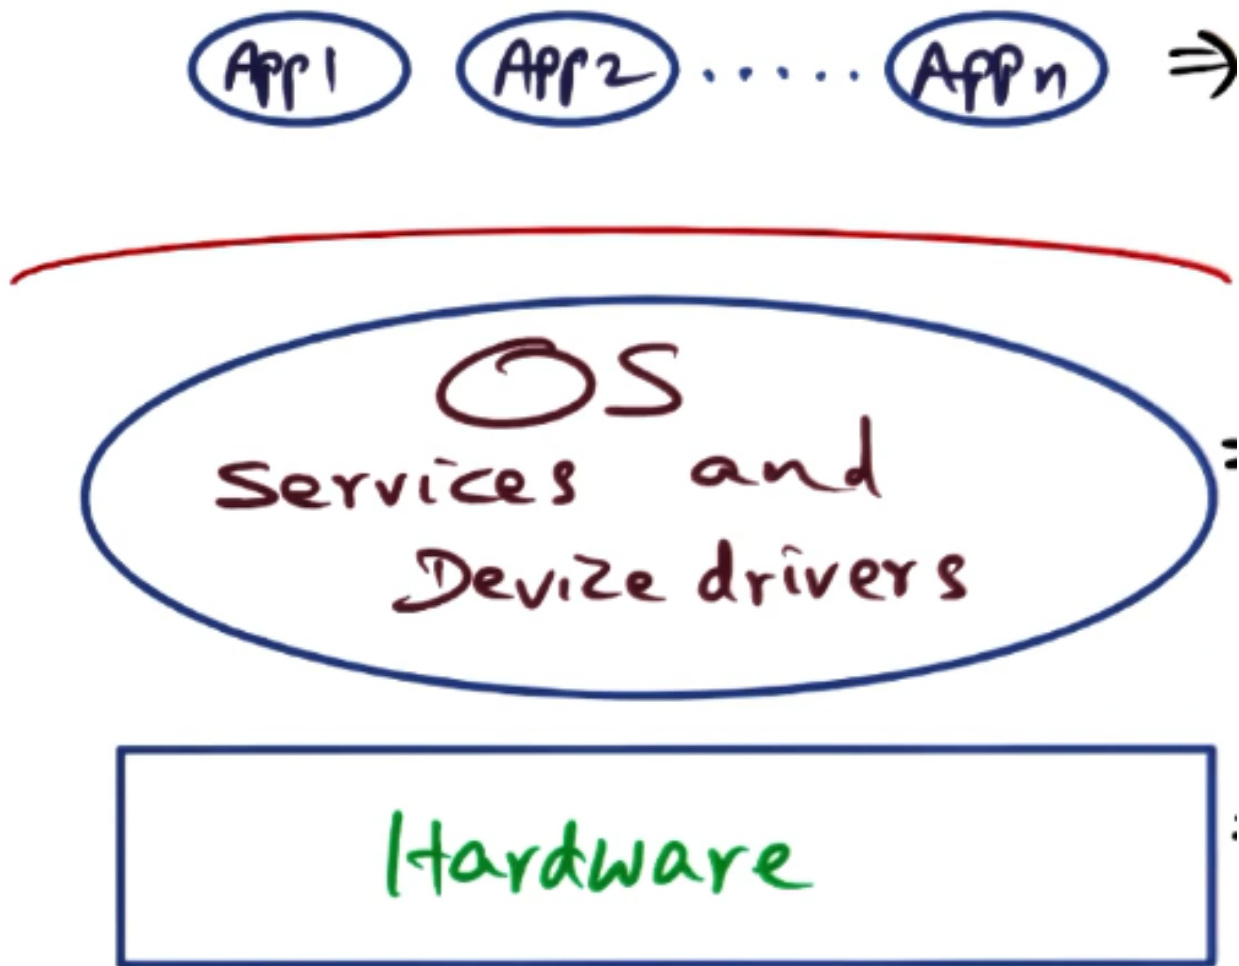# OS Structure

OS Structure Overview

* Operating System Services
    1. Process/thread management and scheduling
    2. Memory management (protection, sharing, demand paging)
    3. Interprocess communication (IPC)
    4. File system
    5. Access to I/O devices such as microphones and speakers
    6. Access to the network
* OS Structure: The way the operating system software is organized with respect to the applications that it serves and the underlying hardware that it manages
* Goals of OS Structure
    1. Protection: Within and across users and to the OS itself
    2. Performance: Time taken to perform the services
    3. Flexibility: Extensibility (not one size fits all)
    4. Scalability: Performance improves with more hardware resources
    5. Agility: Adapting to changes in application needs and/or resource availability
    6. Responsiveness: Reacting to external events
* Commercial OS: Don't meet all goals laid out above
    - Linux
    - MacOS
    - Windows
* Monolithic Structure
    - Each app has its own hardware address space
    - OS also has its own hardware address space (protected from apps)
    - Hardware is managed by the OS
    - OS services and device drivers are all contained in one blob
    - Reduces performance loss by consolidation
        + All of the components are aware of each other and interactions can be optimized
    - Con: Less extensible (no customization, useful for different apps)
        + Video games: Responsiveness is desired
        + Computing prime numbers: Sustained CPU time is desired

Monolithic OS Structure

* DOS-like Structure
    - Disk operating system
    - No protection layer between the application and the OS
    - Pro: Improved performance (access to system services is like a
    procedure call)
    - Con: Decreased protection (an errant application can corrupt the OS)
    - In the early days of computing, only one process could run at a time
        + Simplicity and performance took precedence
        + No overhead for system calls vs developer procedures
   - Loss of proection is unacceptable for a general-purpose OS
* Opportunities for Customization
    - Memory management (page fault)
        1. OS finds a free page frame (page replacement algorithm)
        2. OS updates page table for faulty process
        3. Resume process
    - The page replacement algorithm can't be ideal for every application
    - Can customize how page replacement is handled
* Microkernel Structure
    - Each app in its own hardware address space

- Microkernel runs in priveleged mode and provides simple mechanisms
    + Threads
    + Address space
    + IPC
- OS services are implemented on top of microkernel (same privelege as applications)
- No distinction between OS services and applications in terms of how they run
- Pro: Extensibility; services can be easily replaced
- Con: Performance; in a monolithic structure, filesystem and memory management are all within the kernel. In a microkernel structure, an app might go through the kernel, to the filesystem service, back to the microkernel, and back to the app. All of this is IPC overhead.
- Border crossings
    + User space <-> system space copying
    + Explicit costs in changing context
    + Implicit costs due to cache misses (change in locality)
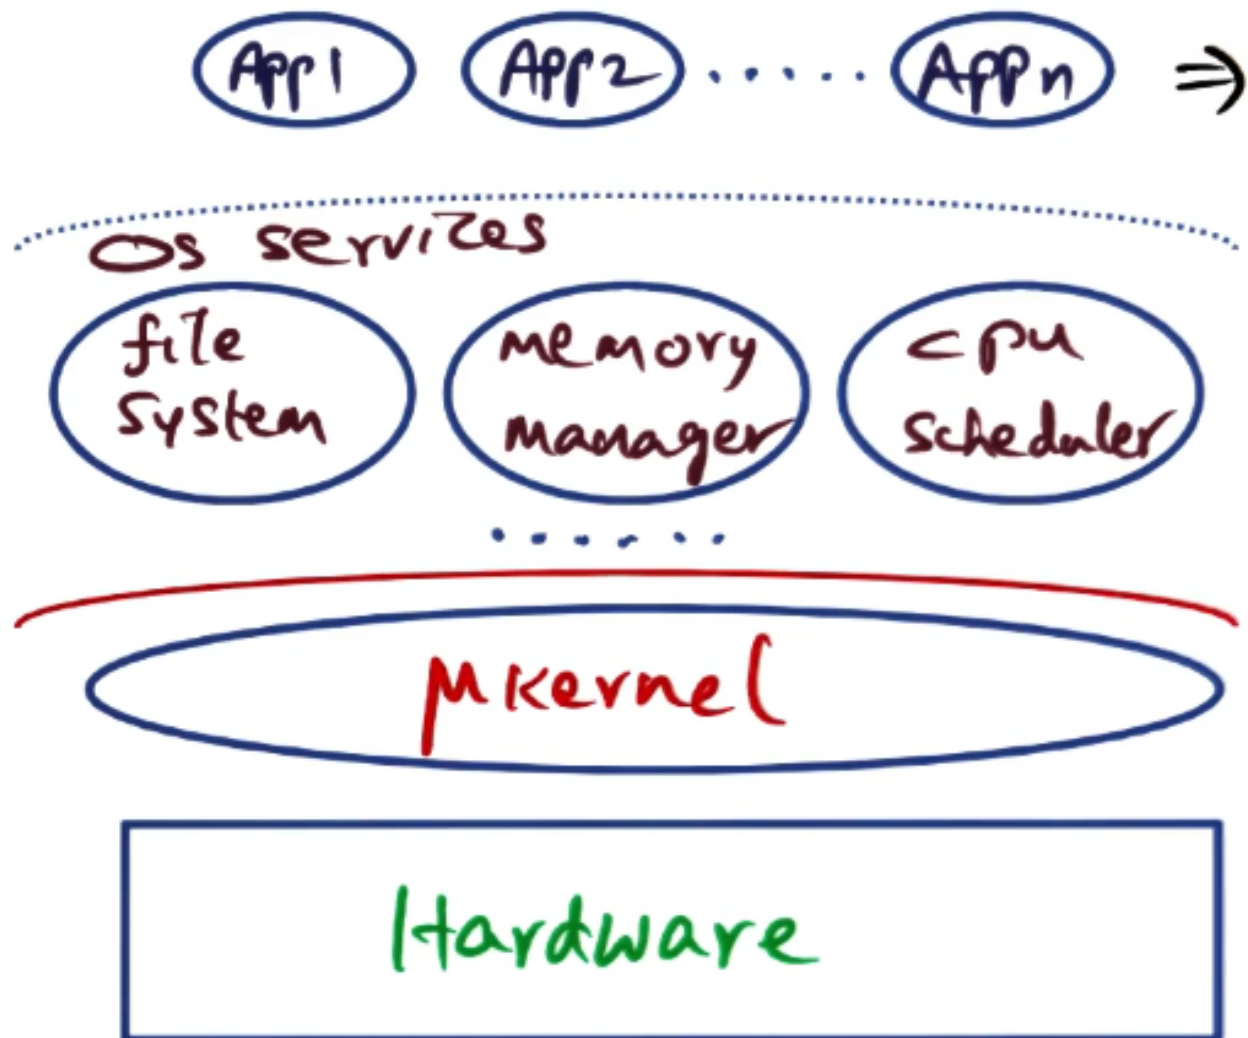
Microkernel OS Structure

Feature	Monolithic	DOS-like	Microkernel
Extensibility		X	X
Protection	X		X
Performance	X	X	

The SPIN Approach

* Two Premises:
    1. Microkernel design compromises on performance due to frequent border crossings
    2. Monolithic design does not lend itself to extensibility
* Goals
    1. Thin (like microkernel) -> only mechanisms, not policies
    2. Access to resources without border crossing (like DOS)
    3. Flexibility for resource management (like microkernel) without sacrificing protection and performance (like monolithic)
* Approaches to Extensibility
    - Hydra OS
        + Kernel mechanisms for resource allocation
        + Capability-based resource access
        + Resource management as coarse-grained objects to reduce border crossing overhead
    - Mach OS
        + Microkernel-based
        + Focused on extensibility and portability
        + Performance wasn't prioritized
* SPIN's Approach to Extensibility
    - Co-location of kernel and extensions to avoid border crossings
    - Compiler enforced modularity (strongly typed language, Modula)
    - Logical protection domains (not hardware address spaces)
    - Dynamic call binding -> flexibility
    - Make extensions as cheap as a procedure call
* Logical Protection Domains
    - Modula-3 safety and encapsulation mechanisms
        + Type safety, automatic storage management
        + Objects, threads, exceptions, generic interfaces
    - Fine-grained protection via capabilities
        + Hardware resources (page frame)
        + Interfaces (page allocation module)
        + Collection of interfaces (entire virtual memory subsystem)
    - Capabilities implemented as language supported pointers
        + Modula-3 pointers are type-specific (no casting)
* SPIN Mechanisms for Protection Domains
    - Create: initialize with object file contents and export names that are contained as entry point methods into the object to be externally visible
    - Resolve: Names between source and target domains
        + Once resolved, resource sharing at memory speeds
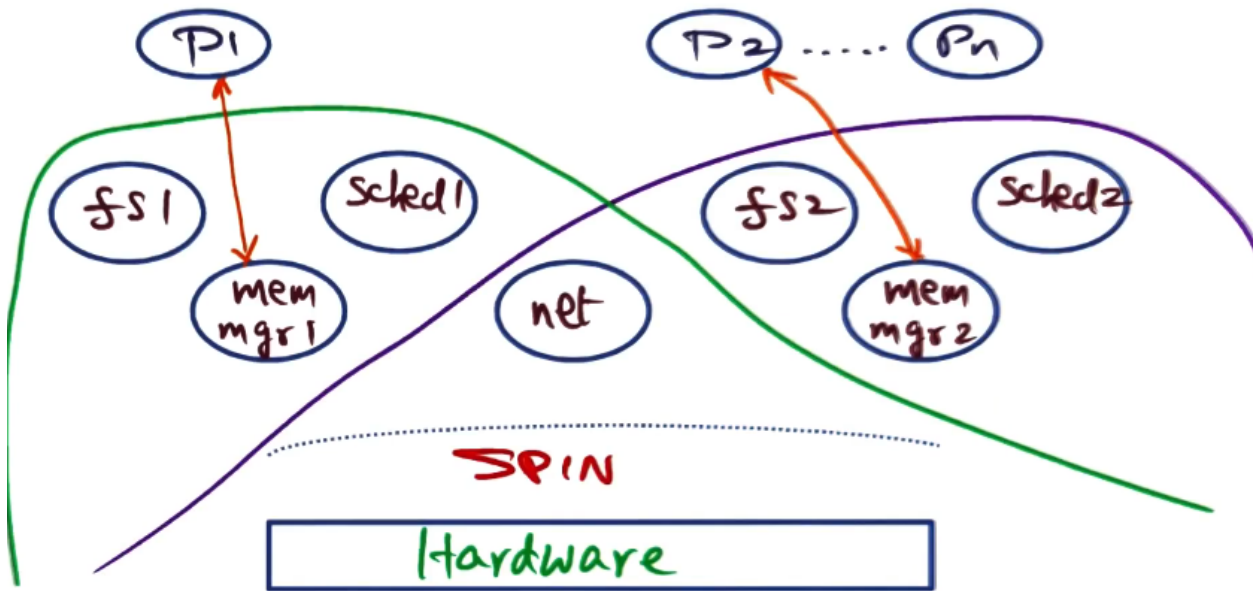   - Combine: To create an aggregate domain from smaller domains
* Customized OS with SPIN
    - No border crossing between services and mechanisms provided by SPIN

Customized OS with SPIN

Create, Resolve, Combine

P1   P2 ..... Pn

fs1   Sched1   fs2   Sched2

mem mgr1   net   mem mgr2

SPIN

Hardware

SPIN OS Structure
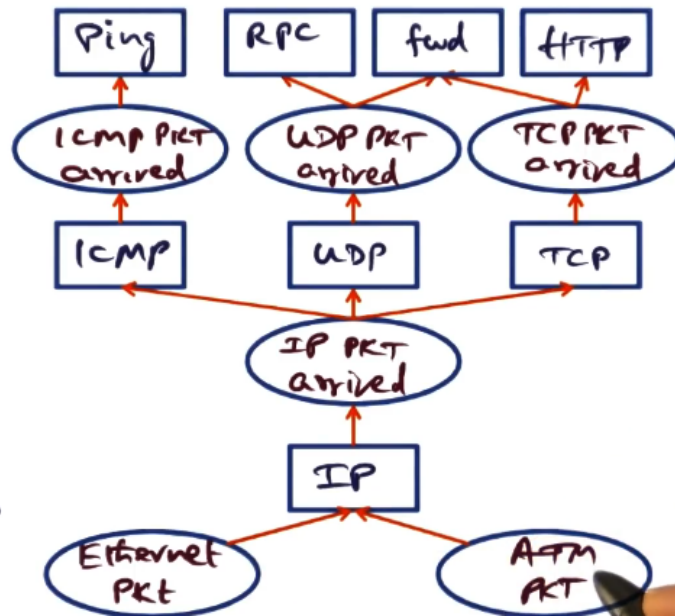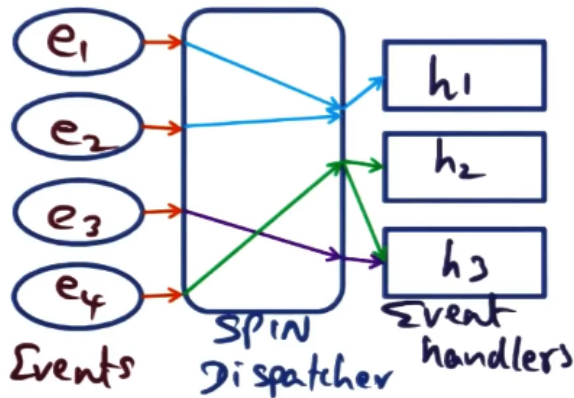
* SPIN Mechanisms for Events
    - OS needs to be able to field events (interrupts, system calls)
    - Event-based communication model
        + Services register event handlers (one-to-one, one-to-many,
        many-to-one)

SPIN Event Handlers

* Default Core Services in SPIN
    - SPIN provides interface procedures for implementing these services
    - Automatically invoked when the hardware event occurs
    - Memory management
        + Physical address (allocate, deallocate, reclaim)
        + Virtual address (allocate, deallocate)
        + Translation (create/destroy address space, add/remove mapping)
        + Event handlers (page fault, access fault, bad address)
    - CPU Scheduling
        + SPIN abstraction: strand (semantics defined by extension)
        + Event handlers (block, unblock, checkpoint, resume)
        + SPIN global scheduler (interacts with application threads package)
    - Core services are trusted because they provide access to hardware
    mechanisms
        + Services may need to step outside the language-enforced protection
        model to control the resources
        + Applications that run on top of extensions must trust extension
        + However, these are isolated (don't impact things that don't rely
   on the extension)

The Exokernel Approach

* Exokernel: Kernel exposes hardware explicitly to OS extensions
    - Decouple authorization of hardware from actual use
    - Expose hardware through secure binding (binds library OS to hardware
    resources)
    - Gives encrypted key to library OS; semantics of use are up to library
        + Can't be forged or transferred
    - Library gives key to exokernel to validate it can use the resource
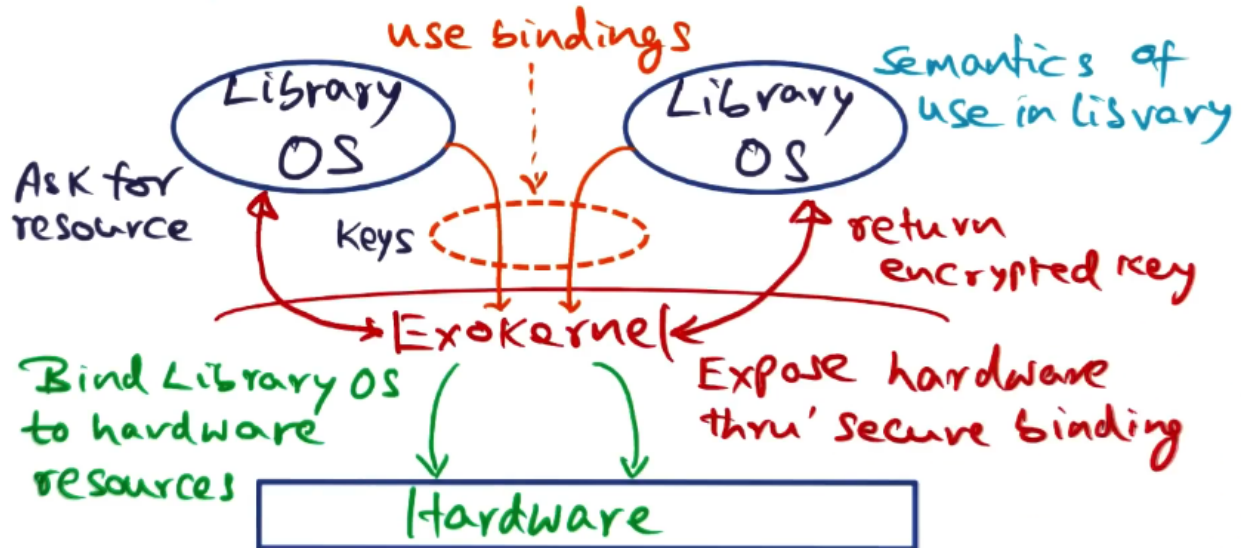
Exokernel OS Structure

* Examples of Candidate Resources
    - TLB Entry
        + Virtual to physical mapping done by library
        + Binding presented to exokernel
        + Exokernel puts it into hardware TLB
        + Process in library OS uses multiple time without exokernel
        intervention
    - Packet filter
        + Predicates loaded into kernel by library OS
        + Check on packet arrival by Exokernel
* Implementing Secure Bindings
    1. Hardware mechanisms (TLB entry, physical page frame, frame buffer)
    2. Software caching
        + "Shadow" TLB in software for each library OS
        + Exokernel dumps hardware TLB into software TLB data structure on
        context switch
    3. Downloading code into kernel
        + Functionally equivalent to SPIN extensions
        + Compromises protection more than SPIN; as long as SPIN's logical
        protection domains are following Modula-3's compile and runtime
        verification, no compromise. Exokernel doesn't make this guarantee
* Memory Management in Exokernel
    - When a thread incurs a page fault, it's caught by Exokernel
    - Then, it's passed to the library OS through the registered handler
    - Library services fault (might require requesting a page frame from
    Exokernel)
* Memory Management Using Software TLB
    - On context switch, much of performance hit comes from loss of

7

cache locality
- To mitigate this overhead, Exokernel implements a software TLB
- Software TLB: Data structure representing a snapshot of the
hardware TLB; captured on context switch for the library OS
- On context switch, Exokernel preloads TLB with software TLB

Default Core Services in Exokernel

Memory Management

Exokernel Memory Management

* CPU Scheduling in Exokernel
    - Linear vector of "time slots"
    - Each library OS will get one time quantum at a time
    - When timing interrupt occurs, Exokernel will switch to next library OS
    - Some time between switching when library OS gets to save state
    - Exokernel only interferes for page faults

Default Core Services in Exokernel

CPU Scheduling
— Linear vector of "time slots"

| OS1 | OS2 | OS1 | OS2 | .... | .... | .... | .... |

T_1 T_2 T_3 T_4

begin end

Time quantum

Exokernel CPU Scheduling

---

* Secure Binding – How can a library OS securely insert code into the kernel?
    - Performance optimization to avoid border crossings
    - In both SPIN and Exokernel, the ability to add extensions must be
    restricted to a trusted set of users
* Revocation of Resources
    - Exokernel needs a way of revoking resources (space/time) that have
    been allocated to a library OS
    - Revoke call is an upcall into library OS (repossession vector)
    - Library OS must clean up when it recevies a repossession vector
        + Stash data in a page to disk when Exokernel reclaims
    - Library can "seed" Exokernel for "autosave"
* Code Usage by Exokernel – Library OS gives code to run on specific events
    1. Packet filter for de-multiplexing of network packets
    2. Run code in Exokernel on behalf of library OS not currently scheduled
    (e.g., garbage collection for an application)
* Summary; Achieving extensibility, protection, and performance
    1. Performance: Exokernel allows for performance critical code in
    library OS to be downloaded securely into Exokernel
    2. Extensibility: Exokernel exposes hardware resources to library OS to
    allow them to implement functionality however is needed
    3. Protection: Exokernel provides encrypted keys to ensure that only the
    intended libary OS can access hardware
    4. Exokernel must be involved when interrupts occur

Exokernel Summary
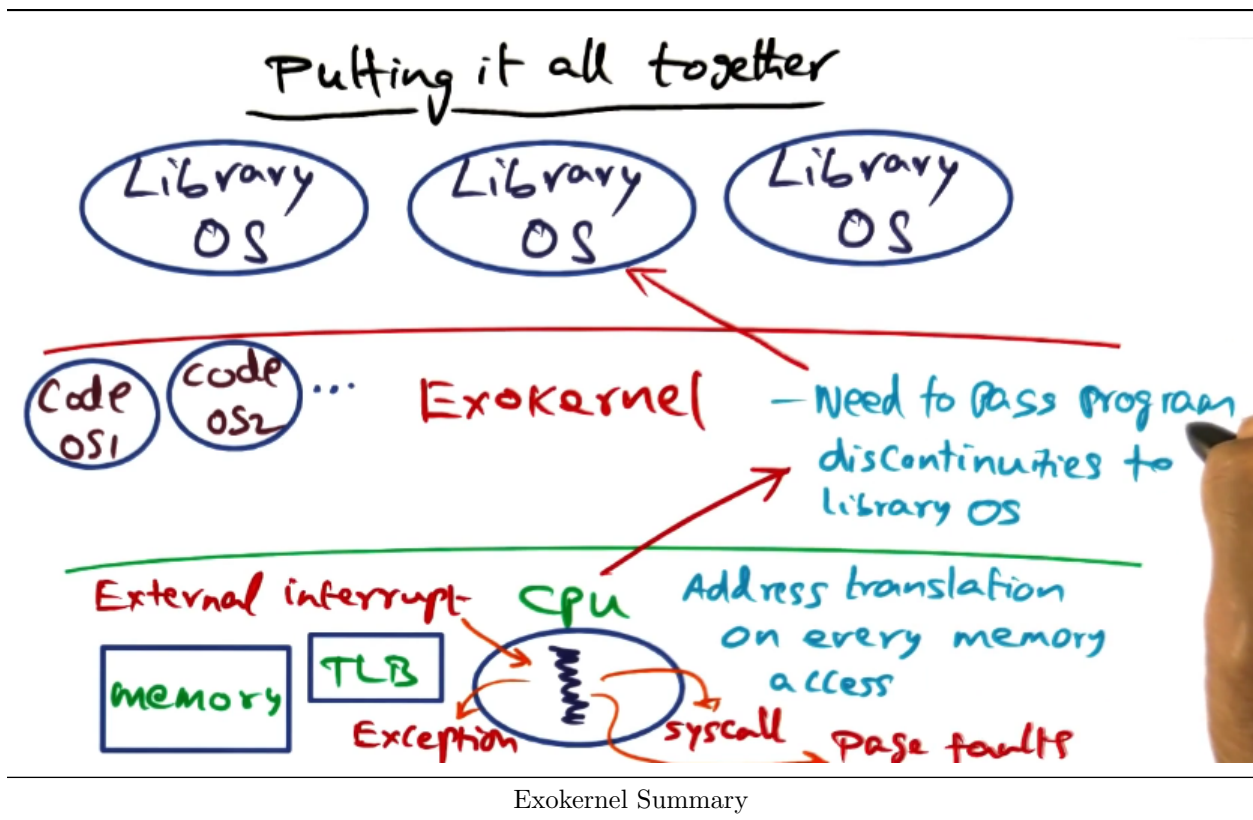
* Exokernel Data Structures
    - PE data structure: Handler entry points for different event types
        + Exceptions, interrupts, system calls, memory mappings
    - Software TLB to preserve page mappings
        + Must specify a set of guaranteed mappings for Exokernel to maintain
        + Not all hardware TLB entries, just the specified ones
    - Similar to event handler structure in SPIN OS
* Performance Results of SPIN and Exokernel
    - Can qualitatively argue that extensibility is provided, but must prove
    quantitatively that performance isn't being sacrificed
    - UNIX: Monolithic
    - Mach: Microkernel (CMU)
    - Both SPIN and Exokernel outperform Mach regarding changing protection
    domains
    - Both SPIN and Exokernel do as well as Unix for dealing with system
    calls

The L3 Microkernel Approach

* SPIN and Exokernel were comparing against CMU's Mach, a microkernel
approach with a focus on portability
   - However, can we design a microkernel with a focus on performance?
* Microkernel: Each service in its own address space
    - Applications sit on top of OS services
    - OS services: Memory management, file system, storage module
    - Microkernel provides simple abstractions (address space, IPC)
* Potentials for Performance Loss
   - Border crossings: Implicit and explicit costs

```
        + Applications and microkernel in different privelege levels
        + System services have to consult other services (cross address space)
     - Protected procedure calls: 100x normal procedure calls
 * L3 Microkernel
     - Aimed to debunk myths about microkernel-based OS structure
     - L3 microkernel provides address space, threads, IPC, and unique IDs
     - L3 puts each service in its own protection domain, not necessarily
     different address spaces
     - All about efficient implementation!
```
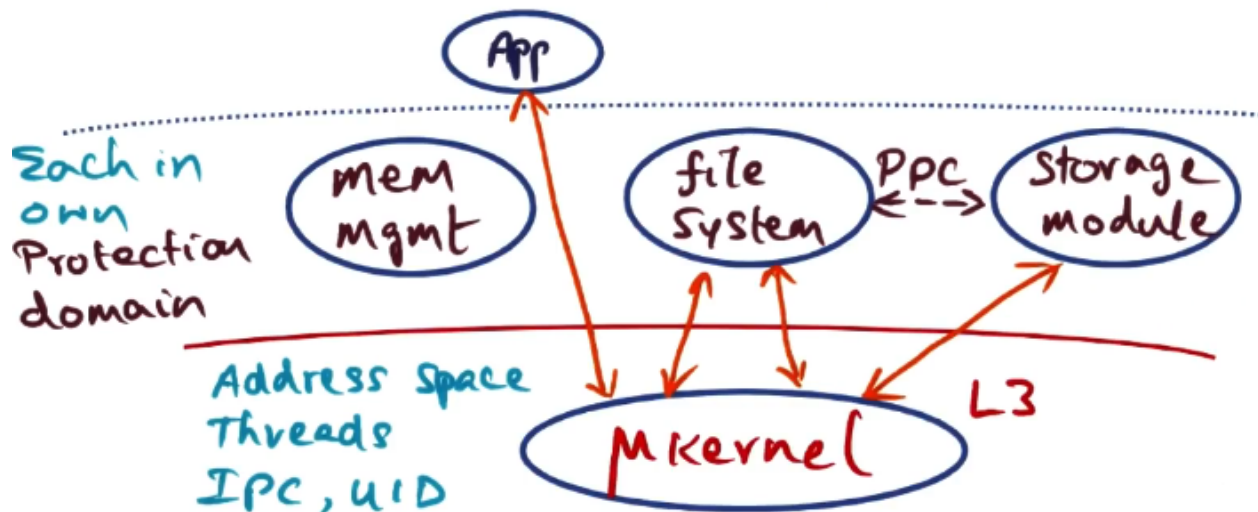

L3 OS Structure

```
 * Strikes Against Microkernel
     - Kernel-user switches (border crossing cost)
     - Address space switches
        + Basis for protected procedure calls for cross protection domain
        calls
        + Crossing address spaces requires flushing the TLB
     - Thread switches and IPC
        + Kernel mediation for protected procedure calls (PPC)
     - Memory effects
        + Locality loss
 * L3 Solution to User Border Crossing
     - L3 accomplishes border crossing in 123 processor cycles
        + Includes TLB and cache misses
     - L3 calculated that the absolute minimum number of cycles on their
     architecture was 107; 123 is pretty good
     - Border crossing isn't inherently slower in their microkernel design
     - For reference, CMU's Mach took 900 cycles on the same architecture
 * L3 Solution to Address Space Switches
     - Address space tagged TLBs: in addition to the tag and index, the TLB
     contains the process for which a particular TLB is present (MIPS)
     - This means the TLB doesn't have to be flushed on context switch
     - Liedtke's recommends exploiting architectural features if TLB is not
     address space-tagged
        + Use segment registers to designate which addresses are valid for
        one running process (segment bounds are hardware enforced); this
```

```
      works well if no process needs the entire address space
         + Shared hardware address space for protection domains
   - Large protection domain: If a process needs the entire address space
   and the hardware doesn't support address-tagged TLBs, a TLB flush is
   required
   - However, this only addresses the explicit cost; loss in cache locality
   is significantly more expensive than the explicit cost
         + Explicit cost << implicit cost
         + TLB flush takes 864 cycles in Pentium processor
   - Small protection domains
         + Switches can be made efficient by careful construction
   - Large protection domains
         + Switching cost not important
         + Cache effects and TLB effects dominate
```
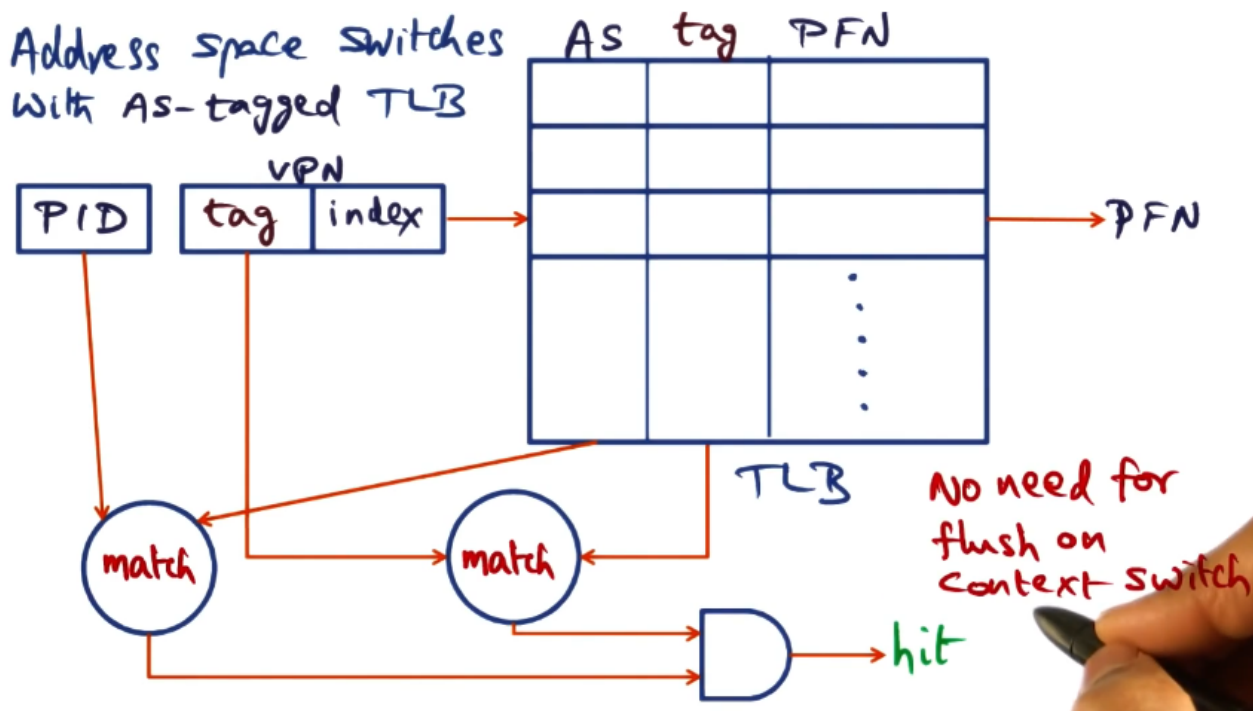

Address Space Tagged TLB

```
* Thread Switches and IPC
   - Shown to be competitive to SPIN and Exokernel by construction
   - Switch involes saving all volatile state of processor
* Memory Effects
   - Assmuption os that loss of locality in a microkernel-based design >>
   monolithic design
   - Liedtke says that splitting the hardware address space into small
   protection domains using segment registers will result in warmer caches
   on context switches
         + Each process only occupies a small memory footprint, which results
         in a smaller footprint in the caches as well
         + Unavoidable if protection domains are large, regardless of how the
         operating system is designed
* Mach's Expensive Border Crossing
```

- Focus on portability (can run on different processor architectures)
            + Code bloat -> large memory footprint -> Less locality -> more
              cache misses -> longer latency for border crossing
        - Kernel memory footprint is the culprit, not microkernel itself
* Thesis of L3 for OS Structuring
    1. Minimal abstractions in microkernel
        + Address spaces, threads, IPC, unique IDs
        + These four abstractions are needed by any subsystem that provides
          functionality to end users
    2. Microkernels are processor-specific in implementation (not portable)
    3. Right set of microkernel abstractions and processor-specific
    implementation -> efficient processor-independent abstractions at higher
    layers
        + Processor-dependent kernel with processor-independent abstractions