

Review

File Systems

1. Key abstractions
 - File
 - Filename (don't need to remember sector number, just name)
 - Directory Tree
2. Access Rights
 - Permissions - Read/Write/Execute for user/group/other
 - Access Control Lists allow for more specific permissions
 - If you don't have execute permissions on a directory, you can't pass through it and modify what's inside
3. Developer's Interface
 - Position/Cursor-based
 - `fd = open("file.txt",...);`
 - `read(fd, ...);`
 - `write(fd, ...);`
 - `lseek(fd, ...);`
 - `close(fd, ...);`
 - Memory-mapped - treat a file as a chunk of memory
 - `fd = open("file.txt",...);`
 - `buf = mmap(..., fd, ...);`
 - `munmap(buf, ...);`
 - `close(fd, ...);`
4. Allocation strategies
 - Platter, track, sector (block)
 - Figures of Merit
 - Simple and fast creation
 - Flexible size
 - Efficient use of space
 - Fast sequential access
 - Fast random access
5. File Allocation Table
 - File is represented as a linked-list of blocks
 - FAT: Indexed by block number
 - Busy: True or False
 - Next: Next block (-1 if end of chain)
 - Directory Table Format (directories treated as files)
 - Filename, starting block, metadata
 - Pros:
 - Easy to create a new file
 - Good for removable storage (copying data; sequential)
 - Spatial efficiency
 - Cons:
 - Bad for random access (tracing pointers)
6. Extended File System
 - Inode structure
 - 12 direct pointers to data
 - 1 indirect pointer to table to data
 - 1 2x indirect pointer
 - 1 3x indirect pointer
 - Directories point to other inode structures
 - Pros:

- Much better random access due to tree structure vs FAT
- Cons:
 - Slightly slower access than FAT due to levels of indirection
- 7. Unified Buffer Cache
 - Cache in main memory to avoid having to go to disk as often
 - Dirty bits allow for the delay of writeback for more opportune times
 - If a file has a short lifetime, writeback might not be necessary
 - fsync/msync to flush cache to disk
- 8. Journaling
 - Sudden power failure means anything in memory will be lost
 - Sequential access thousands of times faster than random
 - Store changes on disk sequentially (journal) then write back later
 - Two writes; more time spent writing
- 9. Direct Memory Access
 - Disk can talk to CPU directly instead of only through memory

Optimization	Reduces Writes	Improves Performance	Improves Recovery
Buffer Cache	X	X	
Journaling		X	X
DMA		X	

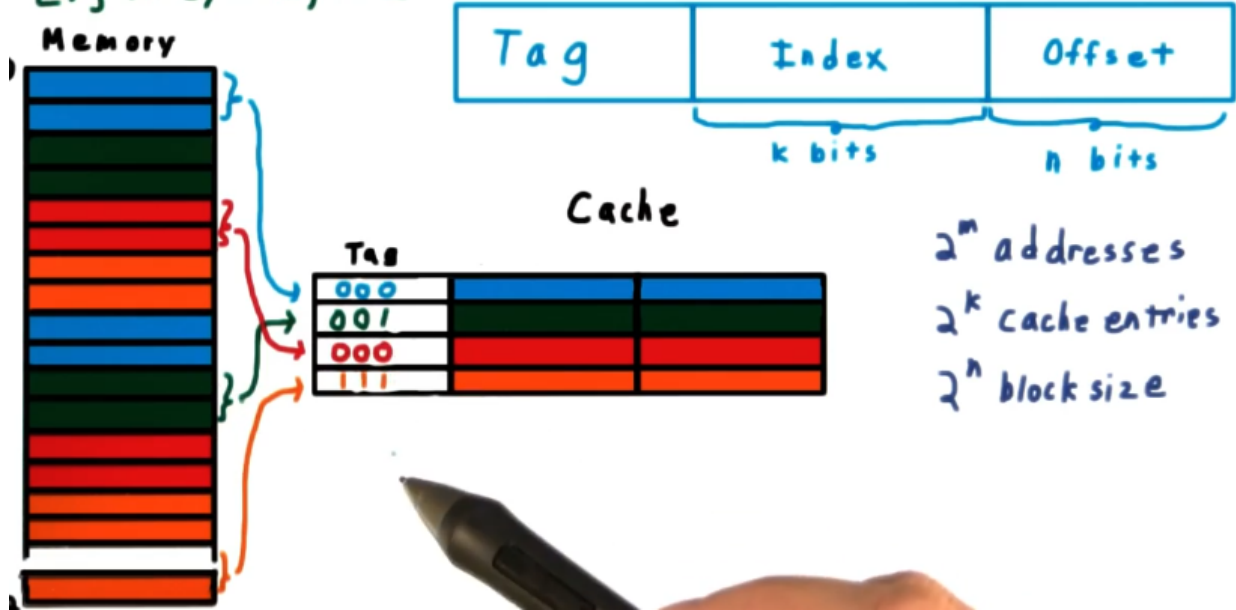
Memory Systems

1. Naive Memory Model
 - Memory
 - Fast
 - Random access
 - Temporary
 - Disk
 - Slow
 - Sequential access
 - Durable
2. Memory Hierarchy (smaller/faster at top of hierarchy)
 - Registers
 - L1 Cache
 - Instruction/data cache
 - L2 Cache
 - L3 Cache
 - Main memory
 - Disk
3. Locality and Cache Blocks
 - Locality: Assume an application will need data near what it already accessed
 - Temporal: Keep data in cache right after it's used
 - Spatial: Instead of a single memory address, fetch entire block
 - 64 bytes, typically
4. Direct Mapping
 - 2^m addresses
 - 2^k cache entries
 - 2^n block size
 - Address: Tag, k bits for index, n bits for offset

Tag	Index	Offset
	k bits	n bits

Direct Mapping

E.g $m=5, k=2, n=1$



Cache Design

5. Set Associative Mapping

- In a direct mapping, the address is only associated with one location in the cache
- In a set-associative mapping, we map an address to multiple locations in the cache
 - Must check two tags to see if there's a cache hit
 - Less likely to encounter the problem of constantly evicting a line
- Introduces replacement policies; typically use least recently used

Tag	Index	Offset
	k-j bits	n bits

6. Fully Associative Mapping

- The address can go in any location in the cache
- There is no index anymore, all tag
 - This means the hardware must check tags for every entry
 - Hardware considerations: Typically 64-256 entries

Tag	Offset
	n bits

7. Write Policy

- Hit:

- Write-through: Write to the cache and memory
 - Write-back: Only write to the cache
- Miss:
 - Write-allocate: Write to cache and memory
 - No-write-allocate: Only write to memory
- 8. Virtual Address Abstraction
 - Each process requires the ability to use the entire system's memory
 - However, processes can't access another processes memory
 - Solution: Each process receives virtual memory addresses
 - The OS maps these virtual addresses to physical locations
- 9. Address Translation
 - Operating system must translate a virtual page number to a physical frame number
 - Virtual pages are typically 4KB
- 10. Page Table Implementation
 - Typically a hierarchical implementation to minimize unused addresses
 - 32-bit VA
 - 32-bit PA
 - 4K pages -> 12-bit offset
 - 10 bits to index into page directory
 - 10 bits to index into page table
 - 12 bits to locate the page in the page table
 - Top-level page table is always in memory (not paged)
- 11. Accelerating Address Translation
 - Translation-lookaside buffer caches VA->PA translations
 - When a context switch occurs, the mappings in the TLB are no longer valid. You can either...
 - Flush TLB on context switch (costly)
 - Store and check address space ID
- 12. Page Table Entries
 - Access control: read/write/execute
 - Valid/present: is the page in memory?
 - Dirty: Has the page been written to? If not, don't write to disk on eviction
 - Control caching
- 13. Page Fault
 - What happens when a process accesses a virtual address that isn't mapped to physical memory?
 - OS page fault handler engages
 - Check if request is valid; if not, seg fault
 - Find a page in memory to store the data
 - If there are no free pages, must evict one (page replacement)
 - Load page into memory from disk
 - Update page table and other data structures
 - Restart the process

```

1  def load(virtual_addr):
2      try:
3          physical_addr = translate_addr(virtual_addr)
4      except page_fault:
5          handle_page_fault(virtual_addr)
6          return load(virtual_addr)
7
8      if is_in_cache(physical_addr):
9          return read_from_cache(physical_addr)
10     else:
11         return read_from_memory(physical_addr)
12
13
14  def translate_addr(virtual_addr):
15      if is_in_tlb(virtual_addr):
16          return read_from_tlb(virtual_addr)
17      elif is_access_violation(virtual_addr):
18          raise access_violation
19      elif is_present(virtual_addr):
20          return read_from_page_table(virtual_addr)
21      else:
22          raise page_fault
23

```

Memory-Management Pseudo-Code

14. Virtually Indexed, Physically Tagged Caches

- Only the page offset bits are important for accessing the cache
- This means we can use the virtual page number to access the TLB and the page offset to access the data cache simultaneously
 - Provides a speedup

Multithreaded Programming

1. Process-Thread Relationship

- Each process has a stack, heap, globals, constants, and code
- Each thread has its own stack, but shares everything else

2. Joinable and Detached Threads

- Detached Threads Termination Cases
 - Any thread makes an exit call, or main reaches the end of its code
 - Our thread returns or calls `pthread_exit`
- Joinable threads won't terminate until another thread joins them
 - pthreads are joinable by default

3. Thread Patterns

- Team model: Pool of worker threads fetching requests from queue
- Dispatcher model: Dispatcher thread passes requests to available worker
- Pipeline model: Break work into subtasks

4. Producer-Consumer Pattern
 - One thread producers work and adds it to a shared buffer
 - Another thread removes work from the queue and completes it
5. Mutex Lock
 - Atomic structure for ensuring threads don't simultaneously access the same resource
 - Supported by hardware
6. Mutex vs Synchronization
 - Mutex: Prevent two threads from accessing the same memory
 - Synchronization: Controlling where two threads are in their execution
7. Conditional Wait Variables
 - Prevents one thread from spinning until another thread finishes
 - Releases the mutex while waiting for a signal
 - Must spin on the condition, not if, to truly ensure the condition is met

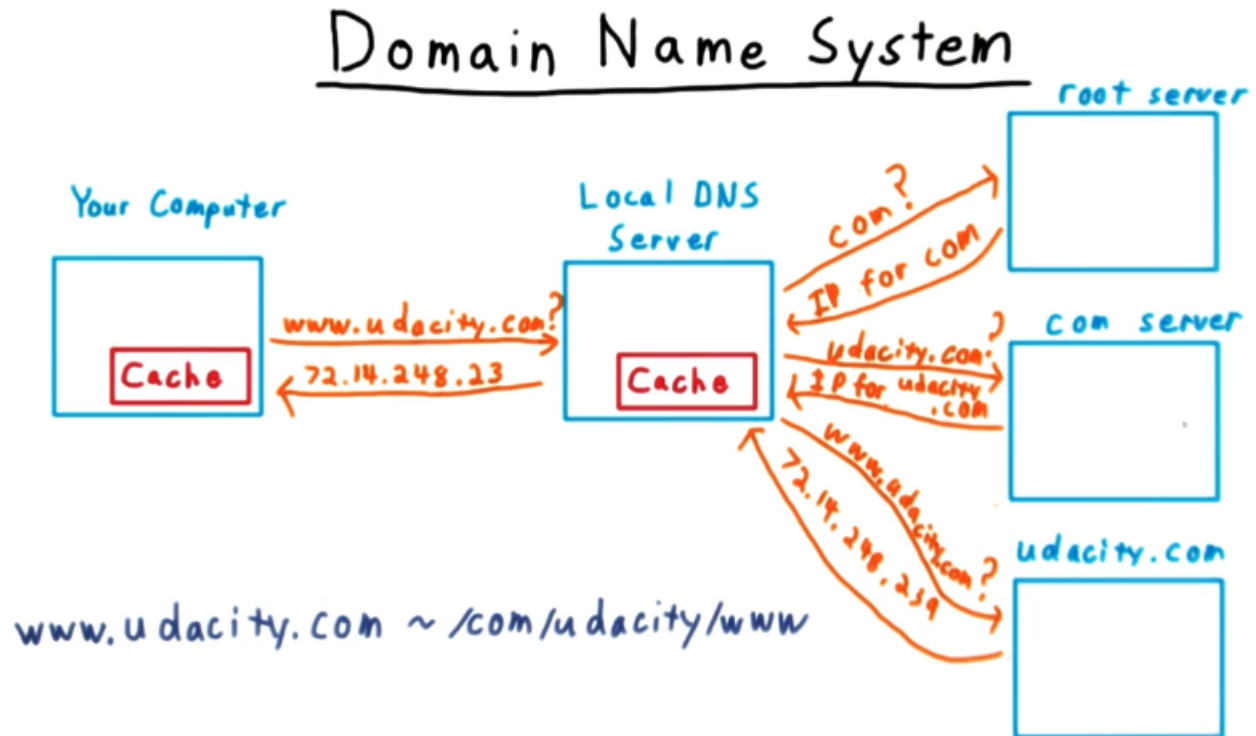
Networking

1. Interconnection Layers
 - Application: Application-specific protocols (HTTP, SMTP, etc.)
 - Transport: Detecting when data is missing, out of order, etc.
 - Network: End-to-end communication on the internet (IP addresses)
 - Link: Communication between peers on a local network (peer-to-peer)
 - Physical: Hardware that creates packets (NIC, modem, Ethernet, etc.)
2. Physical
 - NIC: Direct Memory Access device; doesn't have to go through CPU
 - Only interrupts CPU when data is done being sent or received
3. Link
 - Break large chunks of data into manageable chunks (packets)
 - NICs are identified by a unique 48-bit MAC address (media access control)
 - Every machine on a LAN will receive a message, only the ones that have MAC addresses matching the intended recipient will do anything with it
 - Collision avoidance: Two machines sending messages simultaneously
 - Collision avoidance approaches
 - CSMA/CD: Carrier Sense Multiple Access with Collision Detection
 - Token Ring/Bus
 - Switched Ethernet: Routes frames to destination MAC address; all frames go through the switch

Strategy	Efficient (light loads)	Resilient- Heavy Loads	Fair
CSMA/CD	X		
Token Ring		X	X
Ethernet	X	X	X

4. Network
 - IP address: 32-bit address
 - CIDR Notation: IP/# of bits for network ID
 - v.w.x.y/z -> z represents number of bits that are fixed
 - MIT: 18.0.0.0/8 -> 18.x.y.z are MIT's (2^{24})
 - GT: 130.207.0.0/16 -> 130.207.x.y are GT's (2^{16})
 - IPv6 uses 128-bit IP addresses; IPv4 uses 32-bit
 - Highest in range is broadcast: send to entire subnet
 - Lowest in range refers to subnet as a whole
 - LANS and NAT
 - Local network has a set of private IPs that router knows about

- NAT: Network Address Translation
- ARP: Address Resolution Protocol
- Internet Routing
 - Routing table: Translates IP address to next hop
 - Go through many hops to reach destination
 - netstat -nr
- Autonomous Systems
 - A bunch of discrete nodes that are interconnected
 - BGP: Border Gateway Protocol
- Domain Name System



Domain Name System

5. Transport

- Ports: Specify which process the frame is intended for
 - Allows multiple processes to receive data simultaneously
- Transport layer is only active at the end-points of the graph
- NAT sends the response to whatever port sent it (same for IP)
- Transmission Control Protocol Wishlist
 - Reliability
 - Cope with out-of-order delivery
 - Flow/congestion control
- TCP sends acks and tracks how much data has been sent and received
 - Will only send so many packets without receiving an ack
- User Datagram Protocol
 - Ignores the reliability and out-of-order protection of TCP
 - VoIP: User prefers degradation in quality over delay
 - Streaming: Don't mind occasional packet loss

Contents	UDP	TCP
Port	X	X
Window size		X
Seq/Ack		X

6. Putting it all together

- Application uses DNS to convert host name to IP address
- Routing table points system to modem (can acquire through ARP request)
- Router swaps system port/IP for modem port IP (NAT)
- Modem wraps packet in link layer frame and passes to WAN
- To respond, end user swaps src/dest port and IP and sends packet