

Review

File Systems

- * Key abstractions
 1. File
 2. Filename (don't need to remember sector number, just name)
 3. Directory Tree
- * Access Rights
 - Permissions - Read/Write/Execute for user/group/other
 - Access Control Lists allow for more specific permissions
 - If you don't have execute permissions on a directory, you can't pass through it and modify what's inside
- * Developer's Interface
 - Position/Cursor-based
 - + fd = open("file.txt",...);
 - + read(fd, ...);
 - + write(fd, ...);
 - + lseek(fd, ...);
 - + close(fd, ...);
 - Memory-mapped - treat a file as a chunk of memory
 - + fd = open("file.txt",...);
 - + buf = mmap(..., fd, ...);
 - + munmap(buf, ...);
 - + close(fd, ...);
- * Allocation strategies
 - Platter, track, sector (block)
 - Figures of Merit
 1. Simple and fast creation
 2. Flexible size
 3. Efficient use of space
 4. Fast sequential access
 5. Fast random access
- * File Allocation Table
 - File is represented as a linked-list of blocks
 - FAT: Indexed by block number
 - + Busy: True or False
 - + Next: Next block (-1 if end of chain)
 - Directory Table Format (directories treated as files)
 - + Filename, starting block, metadata
 - Pros:
 1. Easy to create a new file
 2. Good for removable storage (copying data; sequential)
 3. Spatial efficiency
 - Cons:
 1. Bad for random access (tracing pointers)
- * Extended File System
 - Inode structure
 - + 12 direct pointers to data
 - + 1 indirect pointer to table to data
 - + 1 2x indirect pointer
 - + 1 3x indirect pointer
 - Directories point to other inode structures
 - Pros:

- 1. Much better random access due to tree structure vs FAT
- Cons:
 - 1. Slightly slower access than FAT due to levels of indirection
- * Unified Buffer Cache
 - Cache in main memory to avoid having to go to disk as often
 - Dirty bits allow for the delay of writeback for more opportune times
 - + If a file has a short lifetime, writeback might not be necessary
 - fsync/msync to flush cache to disk
- * Journaling
 - Sudden power failure means anything in memory will be lost
 - Sequential access thousands of times faster than random
 - Store changes on disk sequentially (journal) then write back later
 - + Two writes; more time spent writing
- * Direct Memory Access
 - Disk can talk to CPU directly instead of only through memory

Optimization	Reduces Writes	Improves Performance	Improves Recovery
Buffer Cache	X	X	
Journaling		X	X
DMA		X	

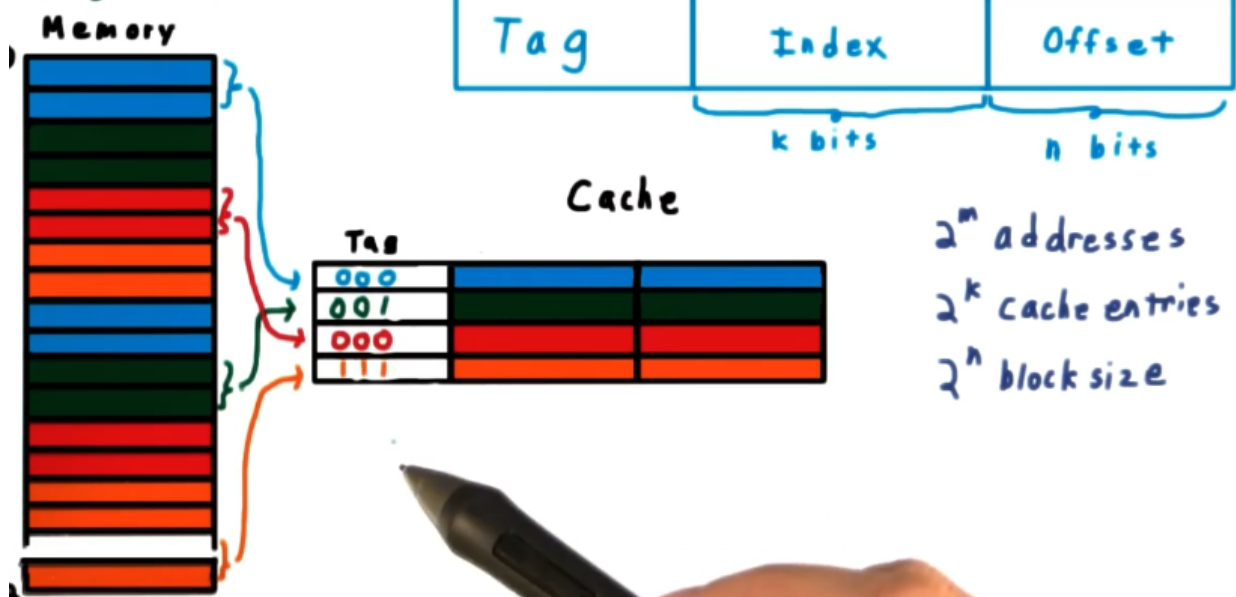
Memory Systems

- * Naive Memory Model
 - Memory
 - 1. Fast
 - 2. Random access
 - 3. Temporary
 - Disk
 - 1. Slow
 - 2. Sequential access
 - 3. Durable
- * Memory Hierarchy (smaller/faster at top of hierarchy)
 - Registers
 - L1 Cache
 - + Instruction/data cache
 - L2 Cache
 - L3 Cache
 - Main memory
 - Disk
- * Locality and Cache Blocks
 - Locality: Assume an application will need data near what it already accessed
 - Temporal: Keep data in cache right after it's used
 - Spatial: Instead of a single memory address, fetch entire block
 - + 64 bytes, typically
- * Direct Mapping
 - 2^m addresses
 - 2^k cache entries
 - 2^n block size
 - Address: Tag, k bits for index, n bits for offset

Tag	Index	Offset
	k bits	n bits

Direct Mapping

E.g $m=5, k=2, n=1$



Cache Design

* Set Associative Mapping

- In a direct mapping, the address is only associated with one location in the cache
- In a set-associative mapping, we map an address to multiple locations in the cache
 - + Must check two tags to see if there's a cache hit
 - + Less likely to encounter the problem of constantly evicting a line
- Introduces replacement policies; typically use least recently used

Tag	Index	Offset
	k-j bits	n bits

* Fully Associative Mapping

- The address can go in any location in the cache
- There is no index anymore, all tag
 - + This means the hardware must check tags for every entry
 - + Hardware considerations: Typically 64-256 entries

Tag	Offset
	n bits

- * Write Policy
 - Hit:
 - + Write-through: Write to the cache and memory
 - + Write-back: Only write to the cache
 - Miss:
 - + Write-allocate: Write to cache and memory
 - + No-write-allocate: Only write to memory
- * Virtual Address Abstraction
 - Each process requires the ability to use the entire system's memory
 - However, processes can't access another processes memory
 - Solution: Each process receives virtual memory addresses
 - + The OS maps these virtual addresses to physical locations
- * Address Translation
 - Operating system must translate a virtual page number to a physical frame number
 - Virtual pages are typically 4KB
- * Page Table Implementation
 - Typically a hierarchical implementation to minimize unused addresses
 - 32-bit VA
 - 32-bit PA
 - 4K pages -> 12-bit offset
 - 10 bits to index into page directory
 - 10 bits to index into page table
 - 12 bits to locate the page in the page table
 - Top-level page table is always in memory (not paged)
- * Accelerating Address Translation
 - Translation-lookaside buffer caches VA->PA translations
 - When a context switch occurs, the mappings in the TLB are no longer valid. You can either...
 - + Flush TLB on context switch (costly)
 - + Store and check address space ID
- * Page Table Entries
 - Access control: read/write/execute
 - Valid/present: is the page in memory?
 - Dirty: Has the page been written to? If not, don't write to disk on eviction
 - Control caching
- * Page Fault
 - What happens when a process accesses a virtual address that isn't mapped to physical memory?
 - OS page fault handler engages
 1. Check if request is valid; if not, seg fault
 2. Find a page in memory to store the data
 3. If there are no free pages, must evict one (page replacement)
 4. Load page into memory from disk
 5. Update page table and other data structures
 6. Restart the process

```

1 def load(virtual_addr):
2     try:
3         physical_addr = translate_addr(virtual_addr)
4     except page_fault:
5         handle_page_fault(virtual_addr)
6         return load(virtual_addr)
7
8     if is_in_cache(physical_addr):
9         return read_from_cache(physical_addr)
10    else:
11        return read_from_memory(physical_addr)
12
13
14 def translate_addr(virtual_addr):
15     if is_in_tlb(virtual_addr):
16         return read_from_tlb(virtual_addr)
17     elif is_access_violation(virtual_addr):
18         raise access_violation
19     elif is_present(virtual_addr):
20         return read_from_page_table(virtual_addr)
21     else:
22         raise page_fault
23

```

Memory-Management Pseudo-Code

- * Virtually Indexed, Physically Tagged Caches
 - Only the page offset bits are important for accessing the cache
 - This means we can use the virtual page number to access the TLB and the page offset to access the data cache simultaneously
 - + Provides a speedup

Multithreaded Programming

- * Process-Thread Relationship
 - Each process has a stack, heap, globals, constants, and code
 - Each thread has its own stack, but shares everything else
- * Joinable and Detached Threads
 - Detached Threads Termination Cases
 1. Any thread makes an exit call, or main reaches the end of its code
 2. Our thread returns or calls pthread_exit
 - Joinable threads won't terminate until another thread joins them
 - + pthreads are joinable by default
- * Thread Patterns
 - Team model: Pool of worker threads fetching requests from queue
 - Dispatcher model: Dispatcher thread passes requests to available worker
 - Pipeline model: Break work into subtasks

- * Producer-Consumer Pattern
 - One thread producers work and adds it to a shared buffer
 - Another thread removes work from the queue and completes it
- * Mutex Lock
 - Atomic structure for ensuring threads don't simultaneously access the same resource
 - Supported by hardware
- * Mutex vs Synchronization
 - Mutex: Prevent two threads from accessing the same memory
 - Synchronization: Controlling where two threads are in their execution
- * Conditional Wait Variables
 - Prevents one thread from spinning until another thread finishes
 - Releases the mutex while waiting for a signal
 - Must spin on the condition, not if, to truly ensure the condition is met

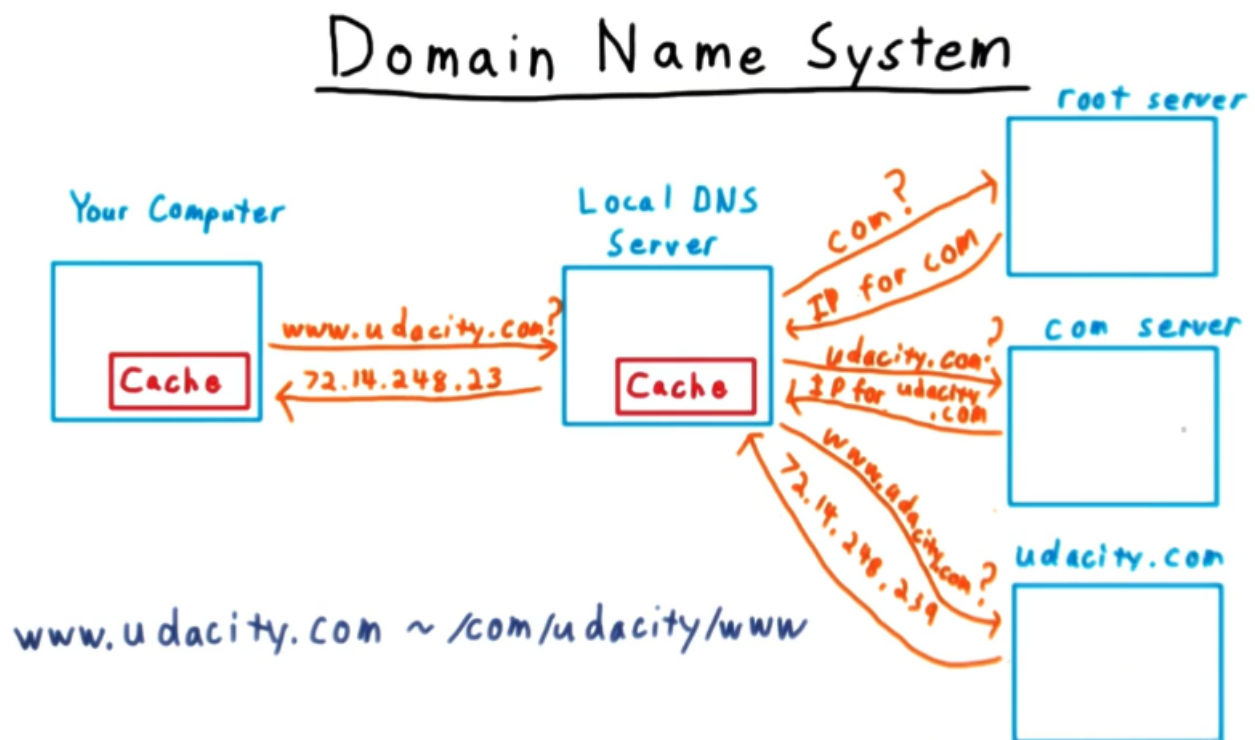
Networking

- * Interconnection Layers
 - Application: Application-specific protocols (HTTP, SMTP, etc.)
 - Transport: Detecting when data is missing, out of order, etc.
 - Network: End-to-end communication on the internet (IP addresses)
 - Link: Communication between peers on a local network (peer-to-peer)
 - Physical: Hardware that creates packets (NIC, modem, Ethernet, etc.)
- * Physical
 - NIC: Direct Memory Access device; doesn't have to go through CPU
 - Only interrupts CPU when data is done being sent or received
- * Link
 - Break large chunks of data into manageable chunks (packets)
 - NICs are identified by a unique 48-bit MAC address (media access control)
 - Every machine on a LAN will receive a message, only the ones that have MAC addresses matching the intended recipient will do anything with it
 - Collision avoidance: Two machines sending messages simultaneously
 - Collision avoidance approaches
 1. CSMA/CD: Carrier Sense Multiple Access with Collision Detection
 2. Token Ring/Bus
 3. Switched Ethernet: Routes frames to destination MAC address; all frames go through the switch

Strategy	Efficient (light loads)	Resilient- Heavy Loads	Fair
CSMA/CD	X		
Token Ring		X	X
Ethernet	X	X	X

- * Network
 - IP address: 32-bit address
 - CIDR Notation: IP/# of bits for network ID
 - + v.w.x.y/z -> z represents number of bits that are fixed
 - + MIT: 18.0.0.0/8 -> 18.x.y.z are MIT's (2^{24})
 - + GT: 130.207.0.0/16 -> 130.207.x.y are GT's (2^{16})
 - IPv6 uses 128-bit IP addresses; IPv4 uses 32-bit
 - Highest in range is broadcast: send to entire subnet
 - Lowest in range refers to subnet as a whole
 - LANS and NAT

- + Local network has a set of private IPs that router knows about
- + NAT: Network Address Translation
- + ARP: Address Resolution Protocol
- Internet Routing
 - + Routing table: Translates IP address to next hop
 - + Go through many hops to reach destination
 - + netstat -nr
- Autonomous Systems
 - + A bunch of discrete nodes that are interconnected
 - + BGP: Border Gateway Protocol
- Domain Name System



Domain Name System

* Transport

- Ports: Specify which process the frame is intended for
 - + Allows multiple processes to receive data simultaneously
- Transport layer is only active at the end-points of the graph
- NAT sends the response to whatever port sent it (same for IP)
- Transmission Control Protocol Wishlist
 1. Reliability
 2. Cope with out-of-order delivery
 3. Flow/congestion control
- TCP sends acks and tracks how much data has been sent and received
 - + Will only send so many packets without receiving an ack
- User Datagram Protocol
 - + Ignores the reliability and out-of-order protection of TCP
 - + VoIP: User prefers degradation in quality over delay
 - + Streaming: Don't mind occasional packet loss

Contents	UDP	TCP
Port	X	X
Window size		X
Seq/Ack		X

* Putting it all together

1. Application uses DNS to convert host name to IP address
2. Routing table points system to modem (can acquire through ARP request)
3. Router swaps system port/IP for modem port IP (NAT)
4. Modem wraps packet in link layer frame and passes to WAN
5. To respond, end user swaps src/dest port and IP and sends packet